

# STREAMING SYSTEMS: KAFKA


*Open-source Technologies for Real-Time Data  
Analytics*

*Imre Lendák, PhD, Associate Professor*


# Overview



- Introduction
- Motivation
- Architecture
- Processes
  - Brokers
  - Producers
  - Consumers
- Replication & consistency
- Monitoring & control



# Definition & origins



- **DEF:** Apache **Kafka** is an open-source data stream processing platform
- **Originally developed by:** LinkedIn
- **Initial release:** Jan 2011
- **Current release:** 2.6.0 in August 2020
- **Written in:** Scala & Java
- **License:** Apache License 2.0
- **Author(s):** 9 core committers, plus ~ 20 contributors
- **Website:** <http://kafka.apache.org/>

# Key features



## Guarantees

- Data integrity checks
- At least once delivery
- In order delivery, per partition

## Characteristics

- Very high performance
- Elastically scalable
- Low operational overhead
- Durable, highly available

---


# MOTIVATION



# Motivation

- LinkedIn's motivation for Kafka was:
  - “A unified platform for handling all the real-time data feeds a large company might have.”
- Features
  - High throughput to support **high volume event feeds**.
  - Support real-time processing of these feeds to create **new, derived feeds**.
  - Support large data backlogs to handle periodic ingestion from **offline systems**.
  - Guarantee **fault-tolerance** in the presence of machine failures.

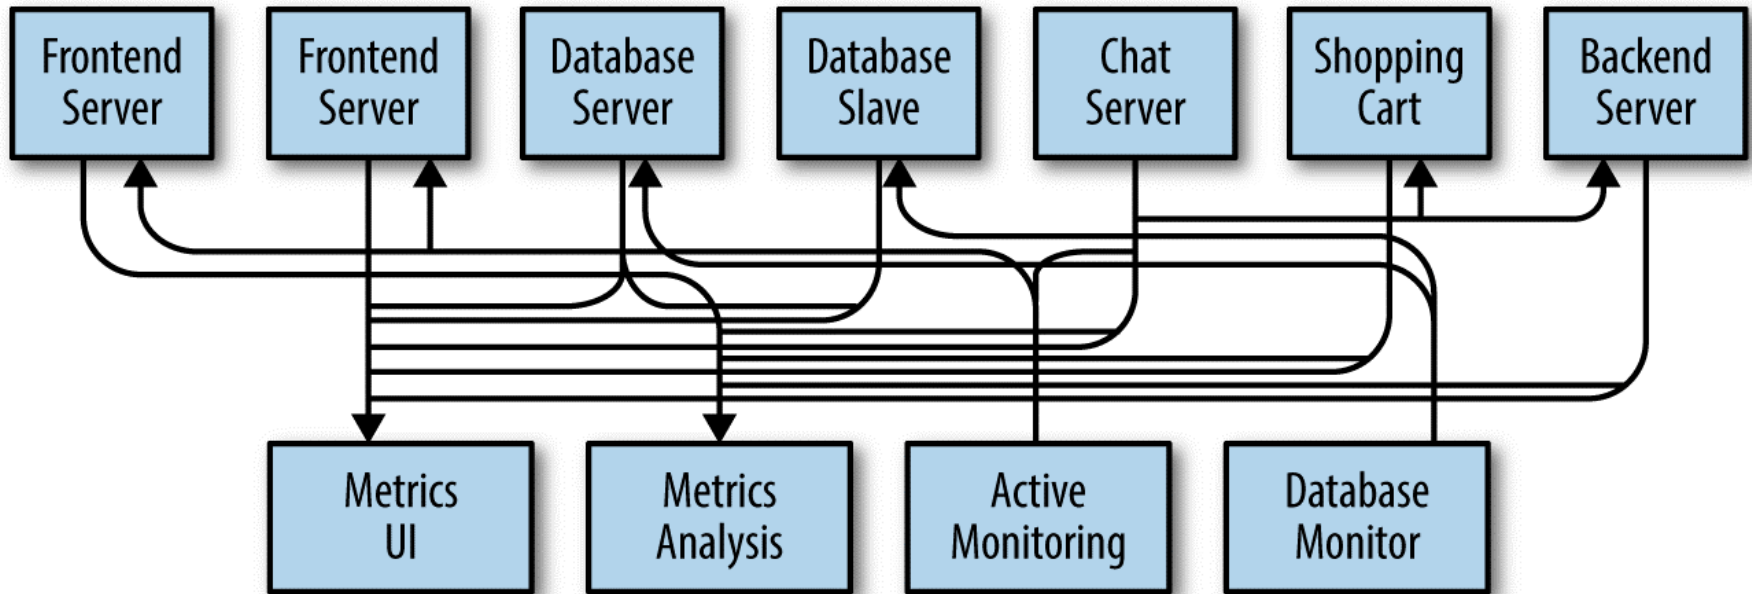
# Scalability challenges – 1



- **Challenge:** What if the number of producers (i.e. server) and consumers (not shown here) increases?

Narkhede, N., Shapira, G., & Palino, T. (2017). *Kafka: the definitive guide: real-time data and stream processing at scale*. " O'Reilly Media, Inc.".

# Scalability challenges – 2

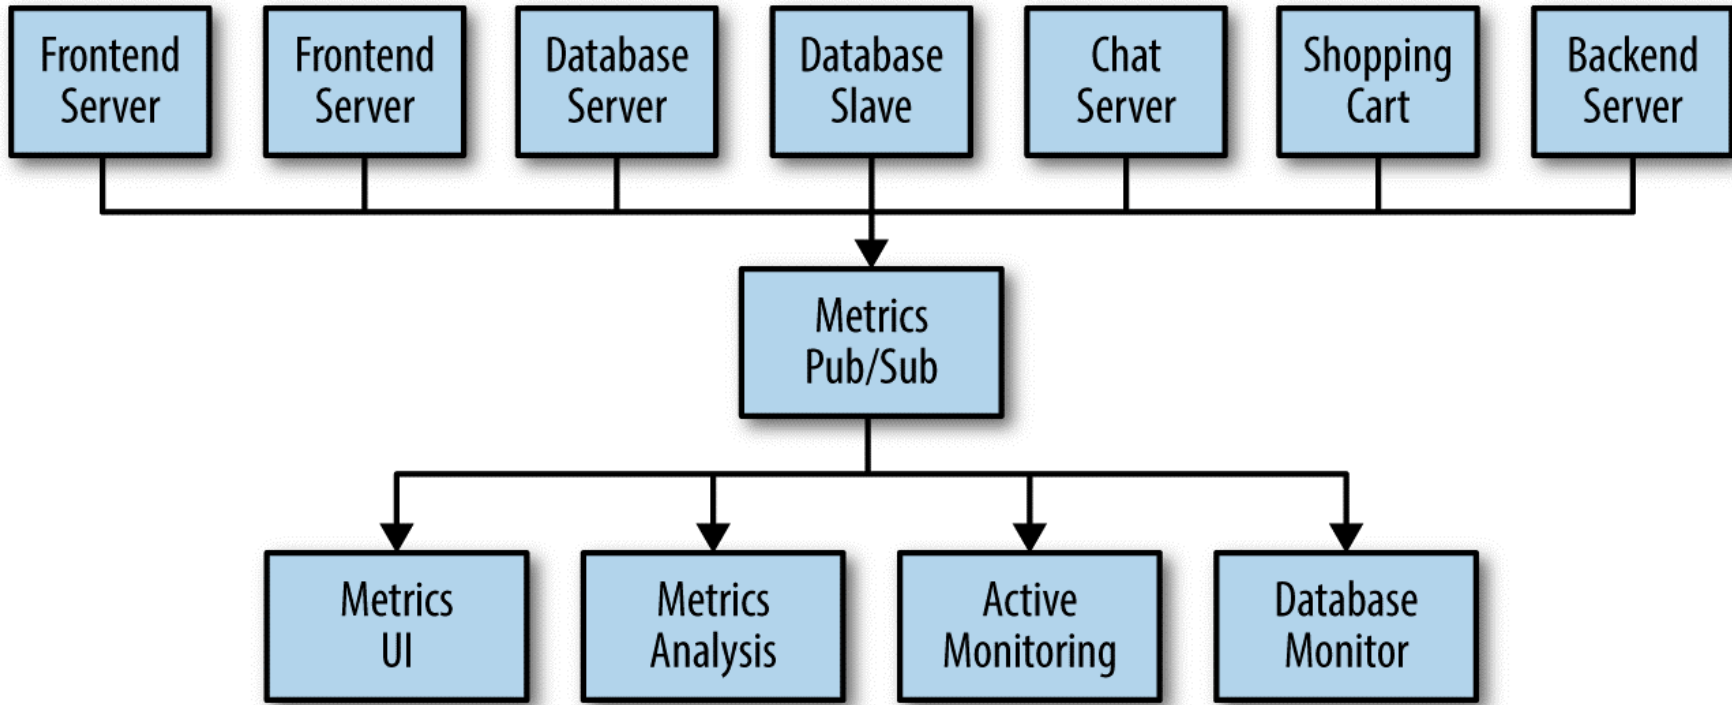


- **Challenge:** How to handle the large number of interconnections between the different sources & sinks?

Narkhede, N., Shapira, G., & Palino, T. (2017). *Kafka: the definitive guide: real-time data and stream processing at scale*. " O'Reilly Media, Inc.".



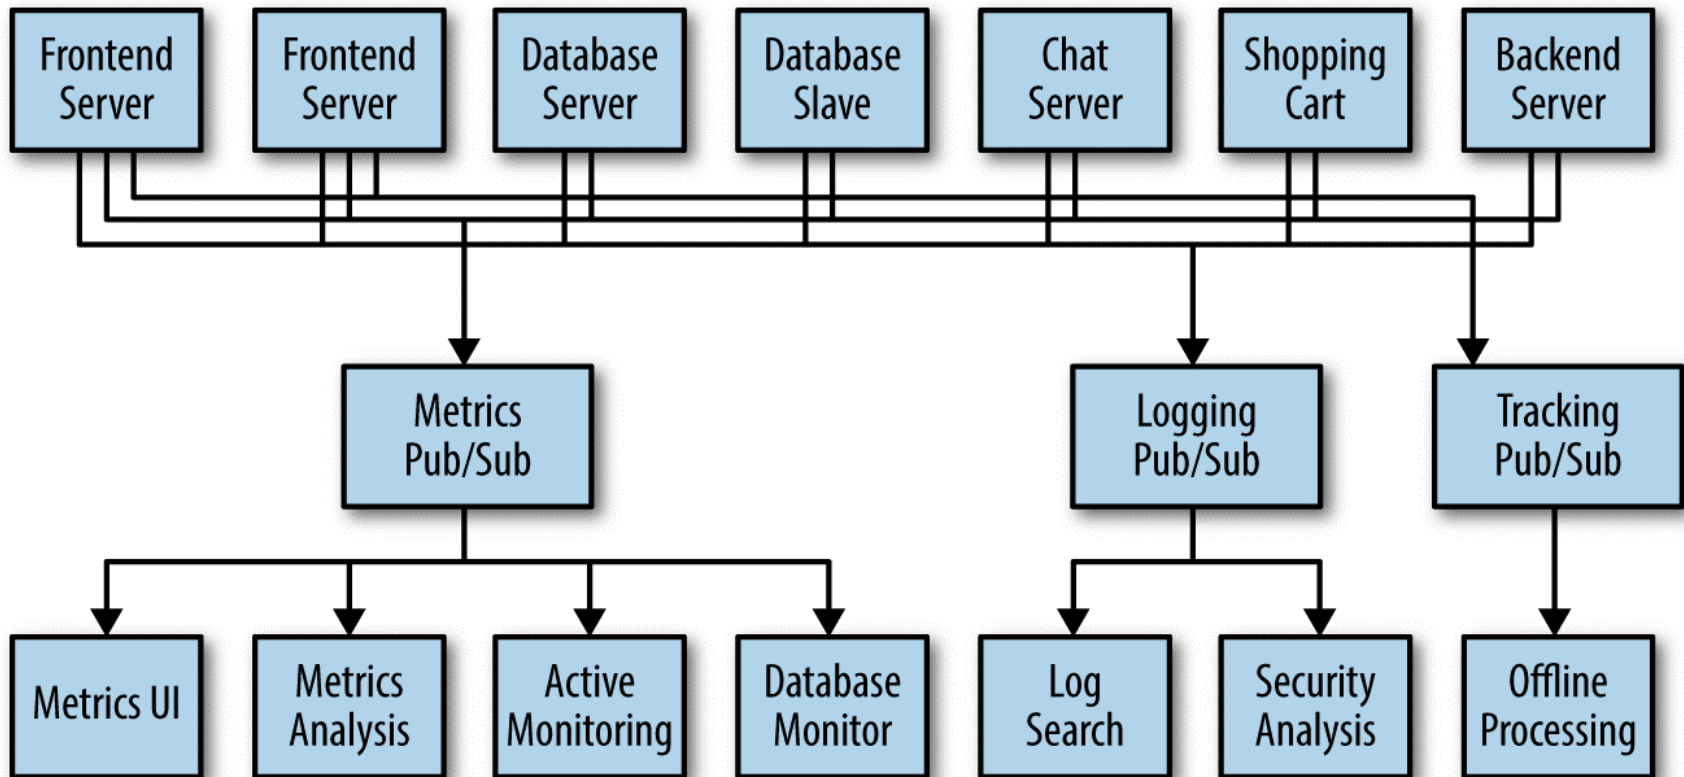
# Scalability challenges – 3



- **Challenge:** Metrics data management solved! But what if there are other data types in a large enterprise?

Narkhede, N., Shapira, G., & Palino, T. (2017). *Kafka: the definitive guide: real-time data and stream processing at scale*. " O'Reilly Media, Inc.".

# Scalability challenges – 4



- **Challenge:** How to handle the diverse data types (in the nodes in the middle?)

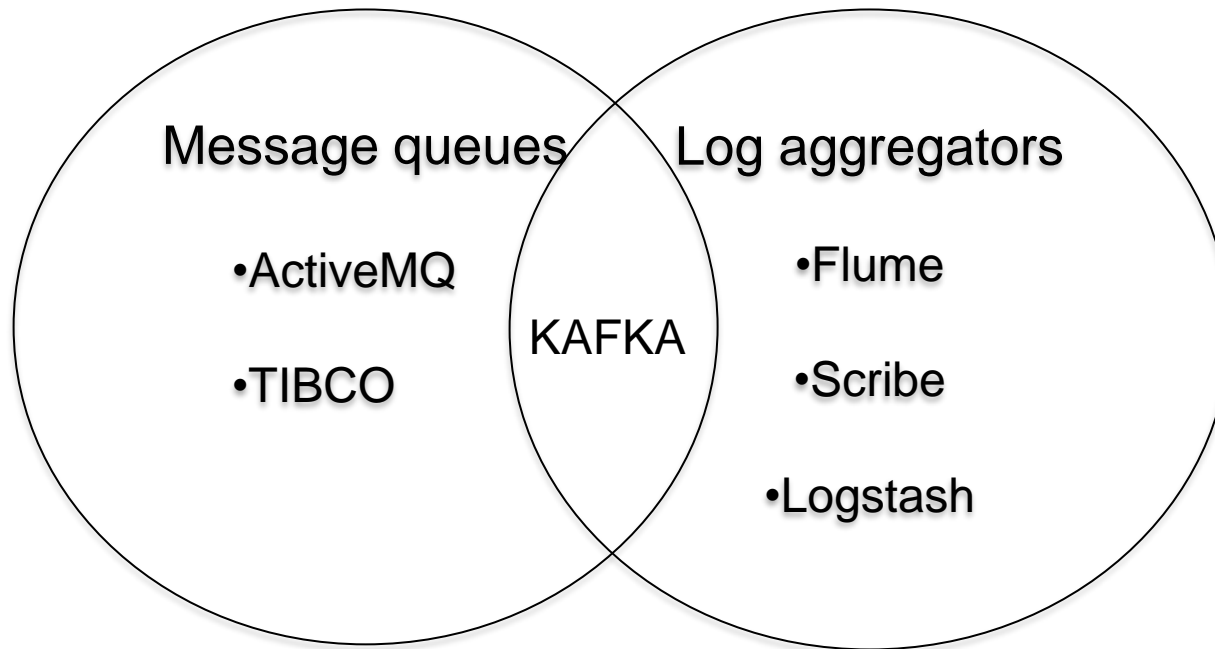
Narkhede, N., Shapira, G., & Palino, T. (2017). *Kafka: the definitive guide: real-time data and stream processing at scale*. " O'Reilly Media, Inc."

---

**ARCHITECTURE**

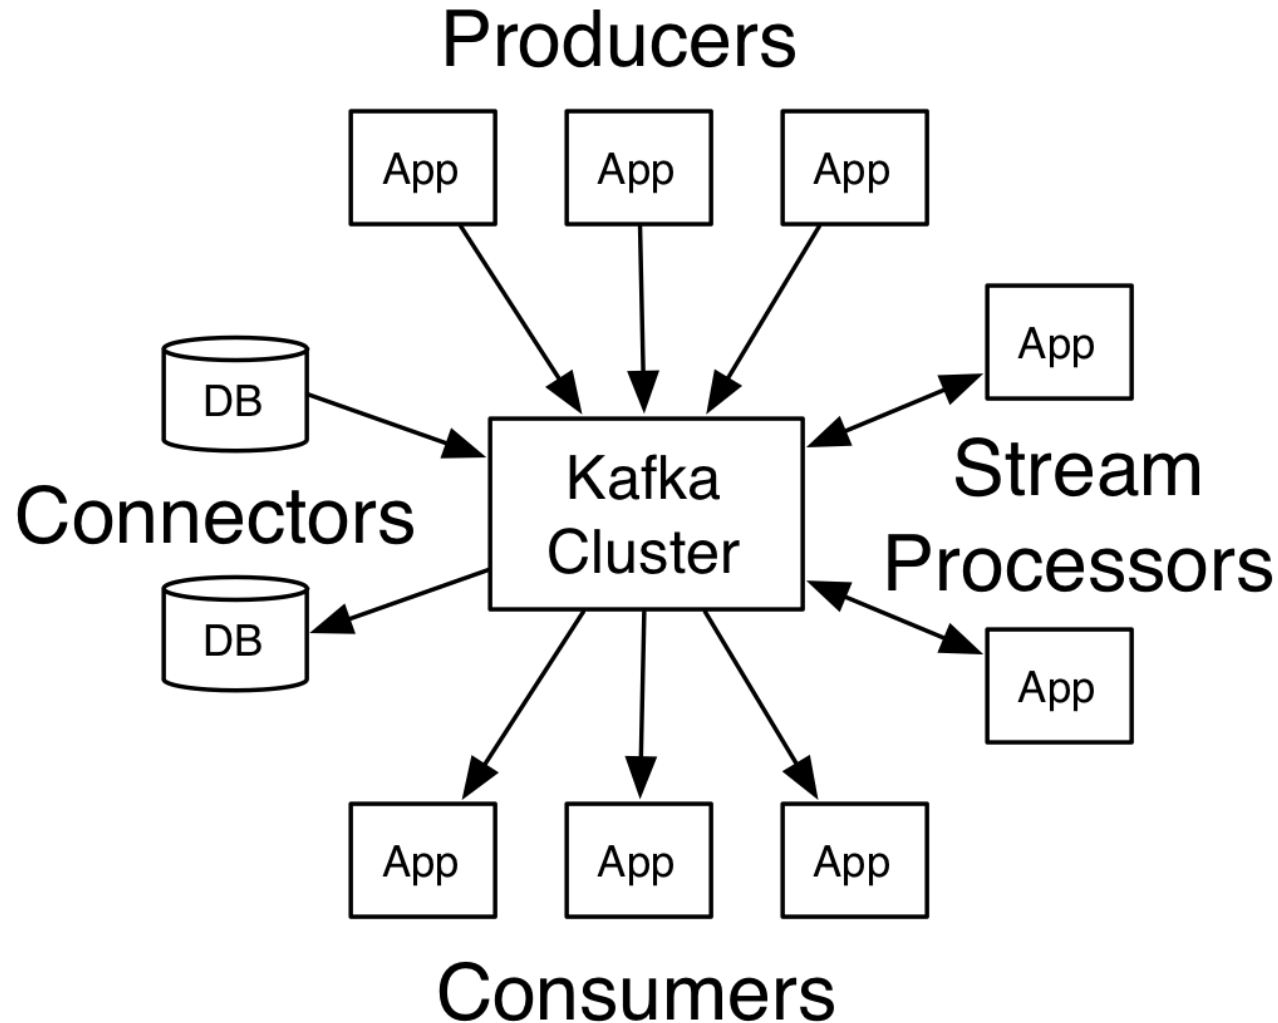


# Queues and aggregators

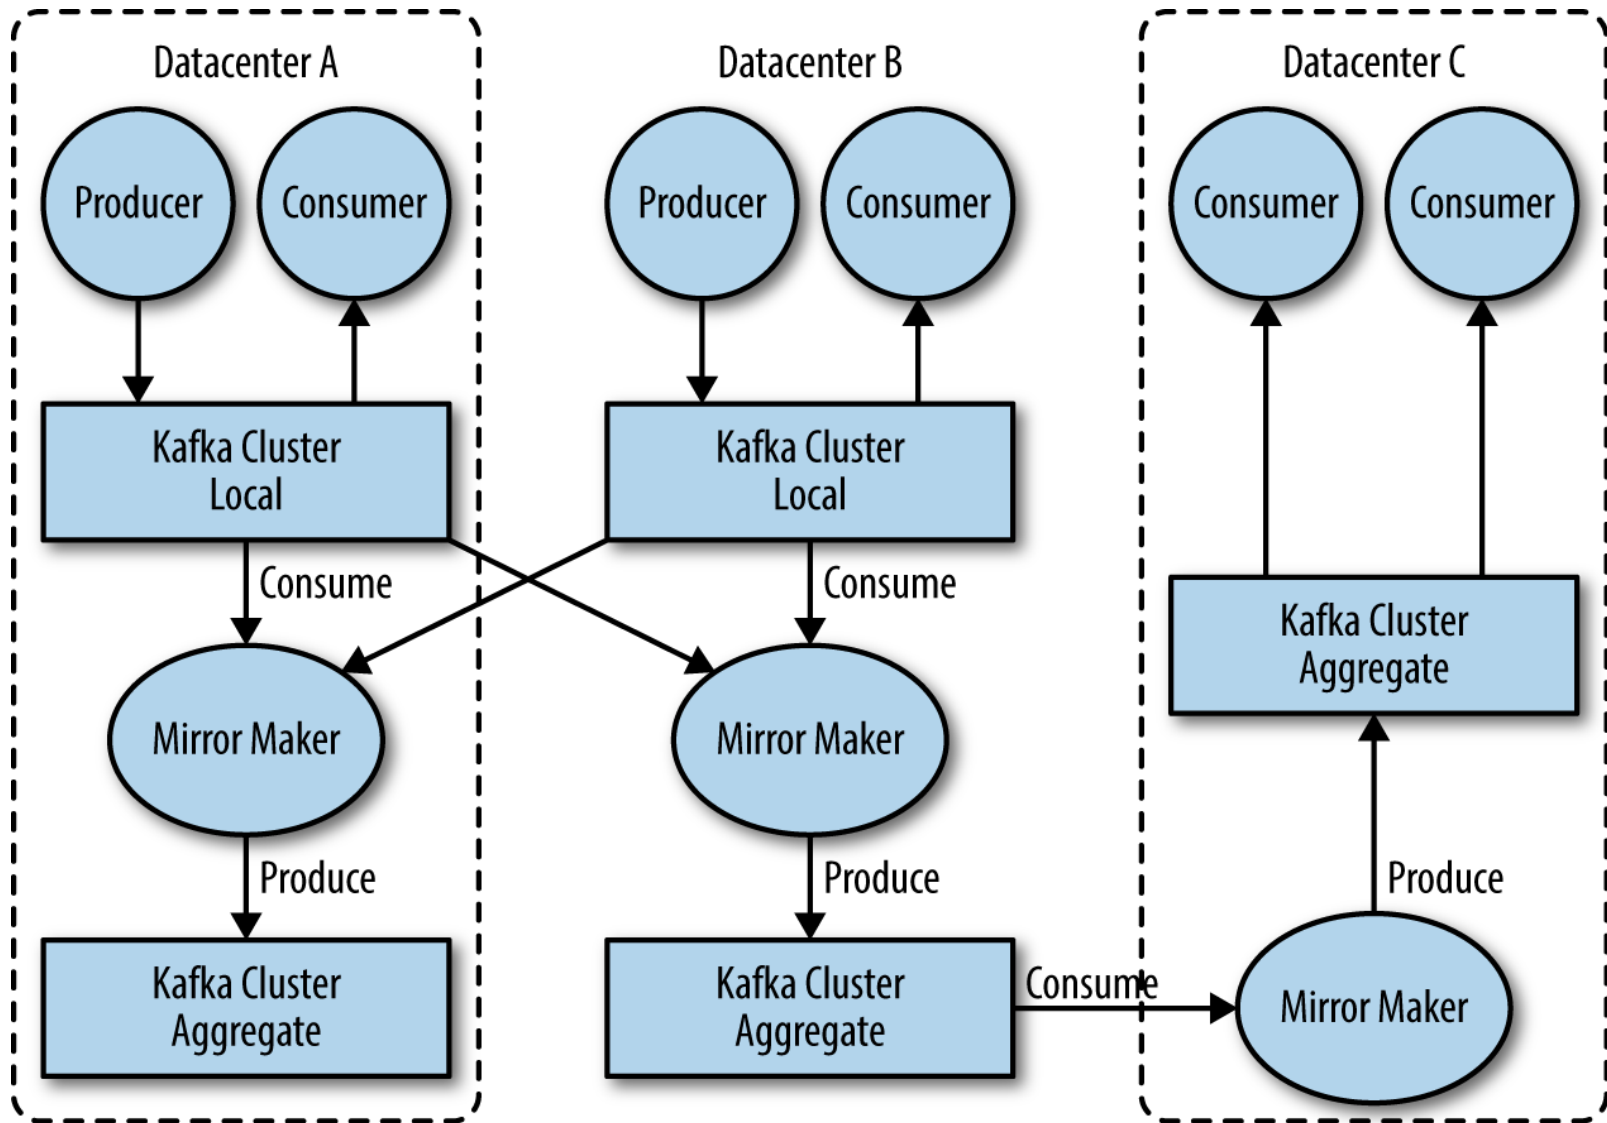


- Kafka has a publish-subscribe (often abbreviated as pubsub) architecture → a distributed mix of message queues and log aggregators

# Architecture overview

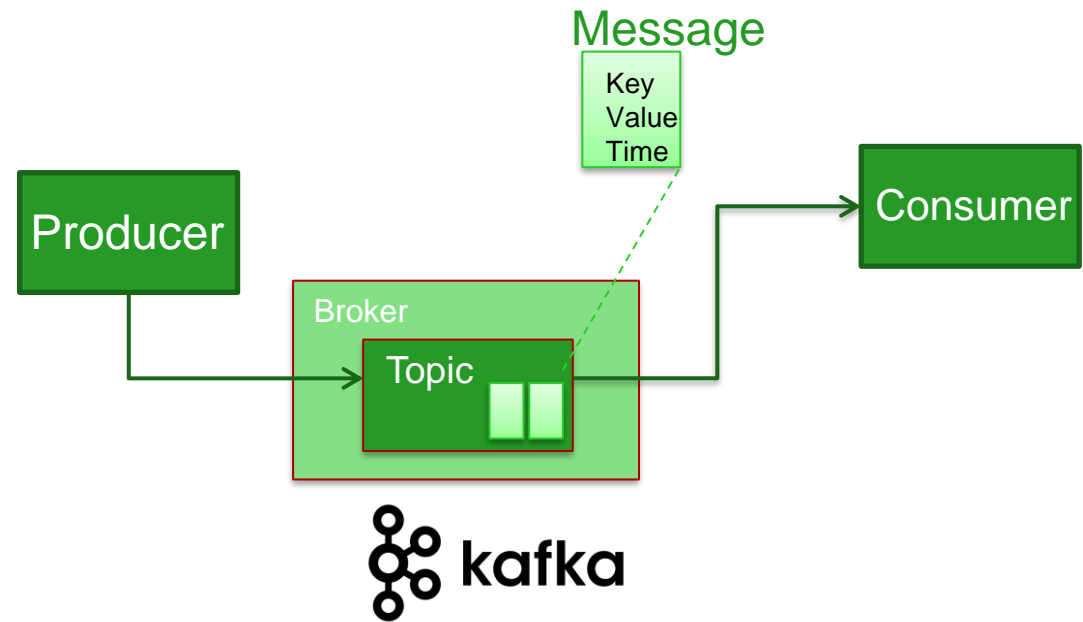


# Multi-datacenter architecture

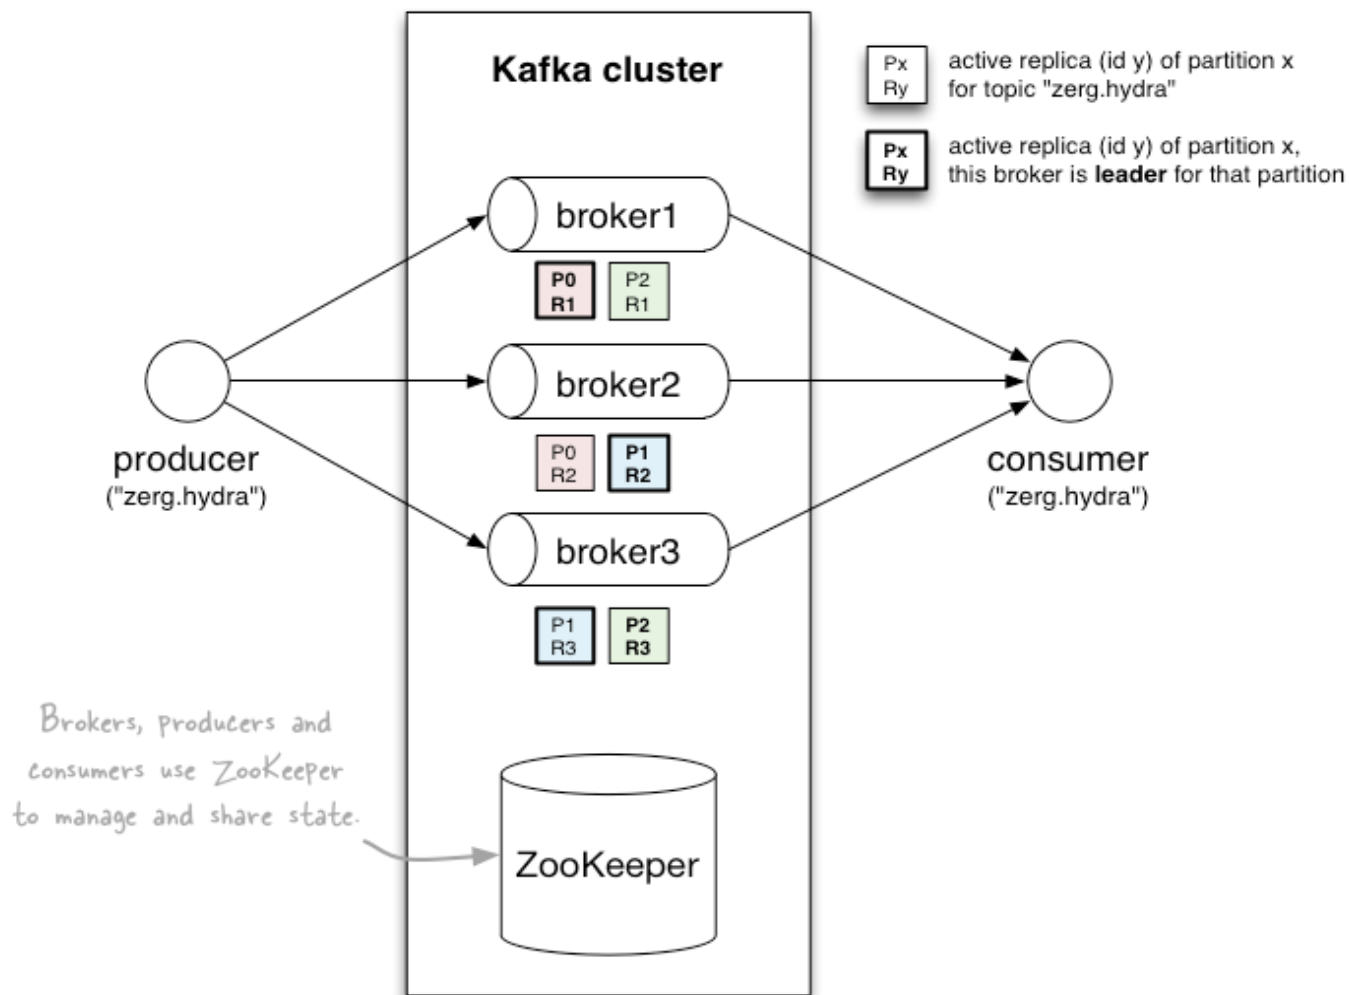


# Who is who? (Terminology)

- Topic
  - partition
- Message
  - == ByteArray
- Broker
  - replicated
- Producer
- Consumer
  - Working together in Consumer Groups



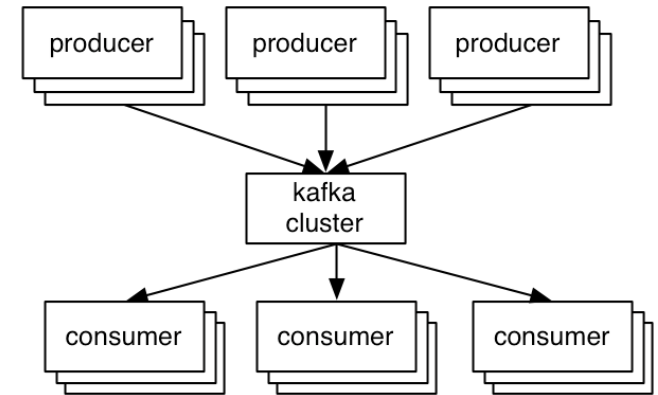
# Architecture





# Key processes

- The who is who
  - **Producers** write data to **brokers**.
  - **Consumers** read data from **brokers**.
- The data
  - Data is stored in **topics**.
  - **Topics** are split into **partitions**, which are **replicated**.

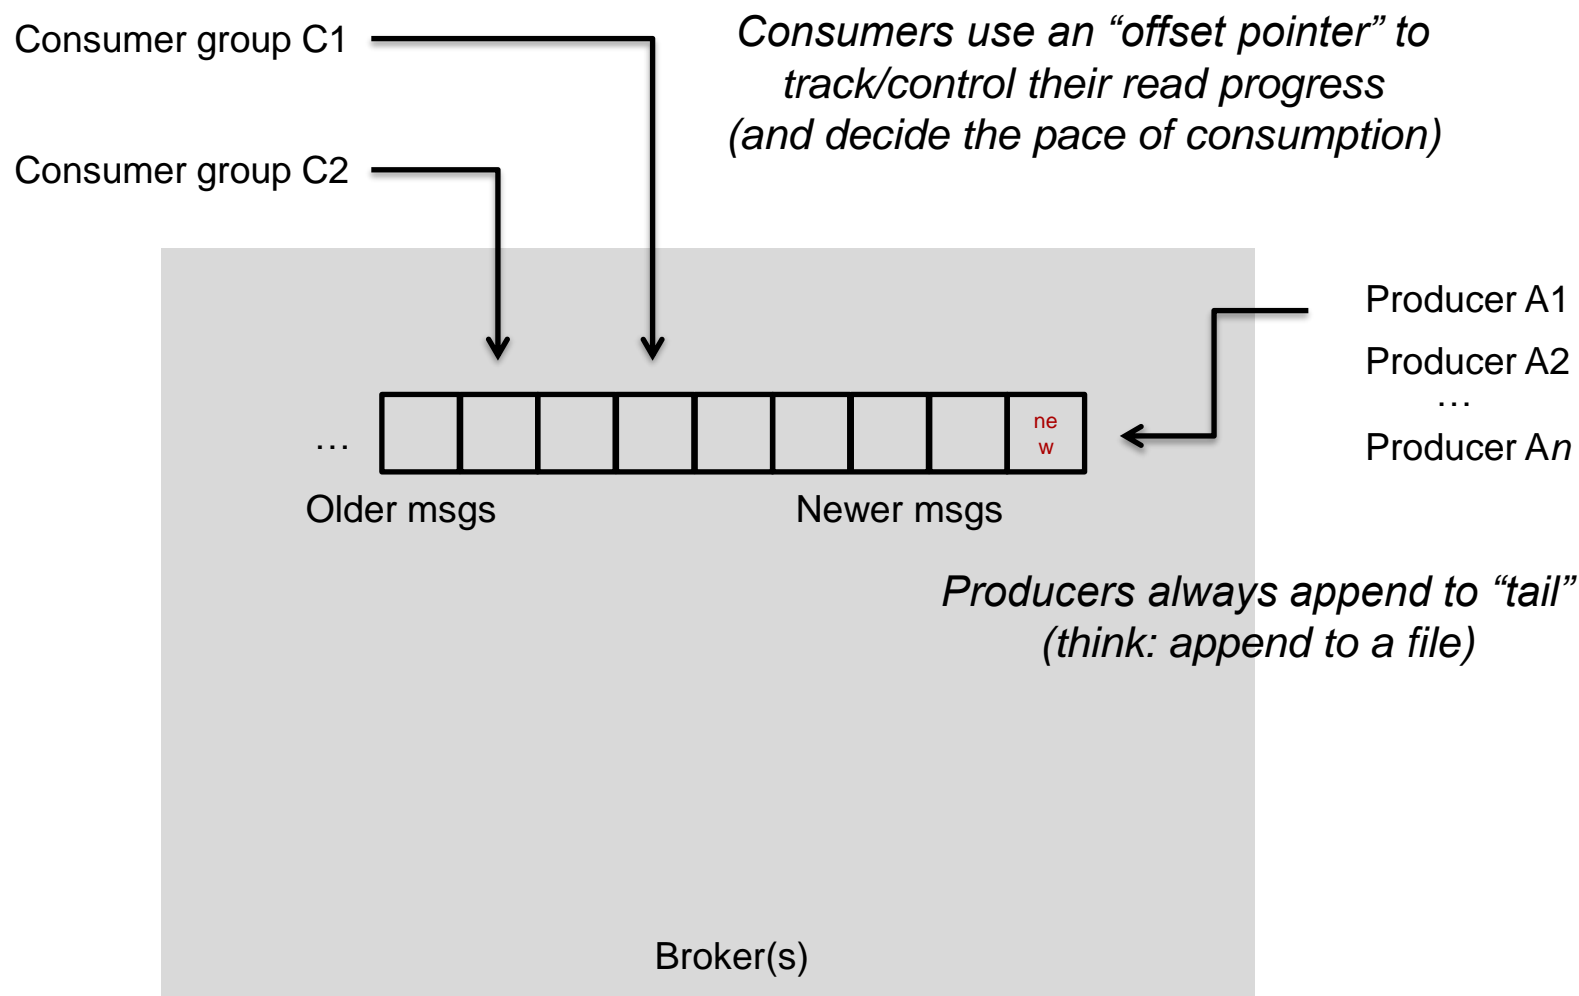


---

**PROCESSES: BROKERS**

- **DEF:** Kafka brokers are the processes tasked to receive messages from producers, consistently store them and respond to requests from Kafka consumers
- A Kafka cluster consists of one or more brokers
- Brokers are usually executed on different servers
- One broker can maintain multiple partitions of different Kafka topics
- The brokers maintain special, non-producer-defined topics for administrative purposes, e.g. topics for memorizing message offsets for consumers

# Topics




# Partitions



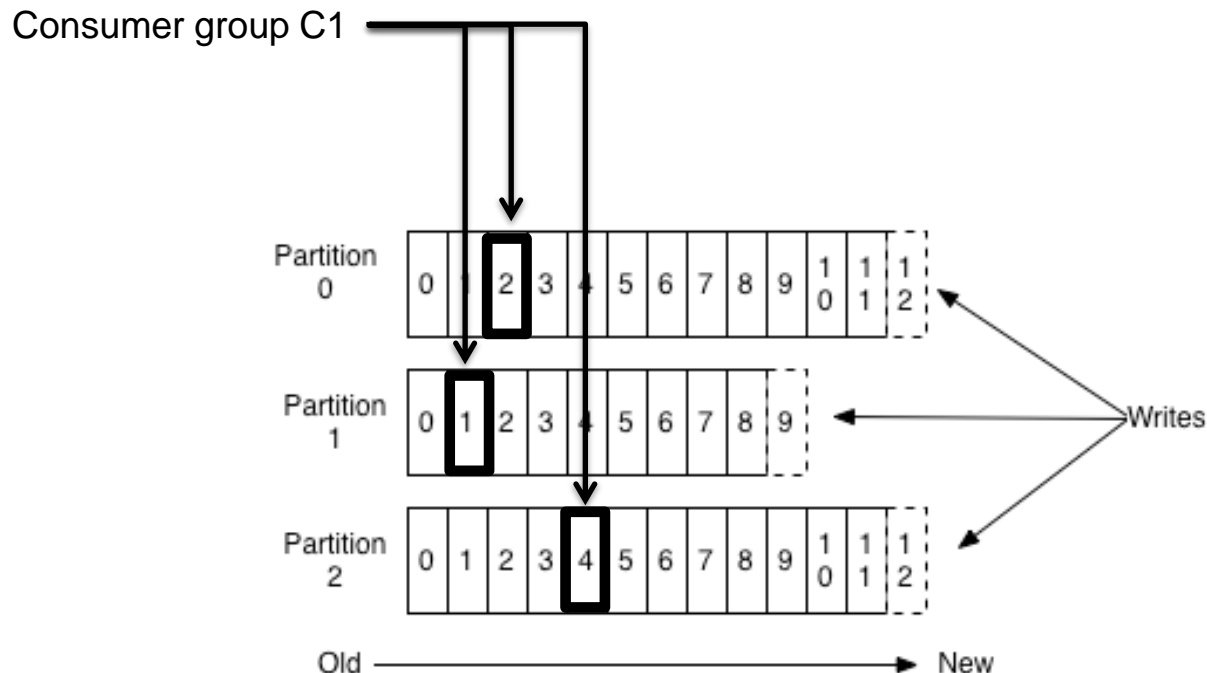
- A topic consists of **partitions**.
- Partition: **ordered + immutable** sequence of messages that is continually appended to

## Anatomy of a Topic



# Partition offsets

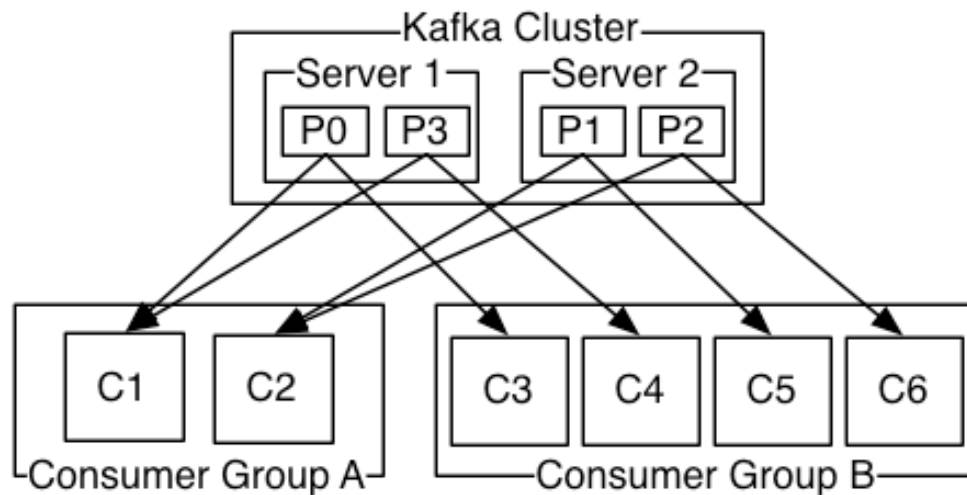
- **Offset:** messages in the partitions are each assigned a unique (per partition) and sequential id called the *offset*
  - Consumers track their pointers via (*offset*, *partition*, *topic*) tuples



# Partitions



- #Partitions of a topic is configurable
- #Partitions determines **max** consumer (group) parallelism



- Consumer group A, with 2 consumers, reads from a 4-partition topic
- Consumer group B, with 4 consumers, reads from the same topic

# Broker performance



## Efficiency

- Each topic has an ever-growing log
  - A log == a list of files
- A message is addressed by a log offset
- Batch send and receive
- No message caching in JVM
- Rely on file system buffering
- 1 file system operation per request

## Implementation

- Fast **writes**:
  - While Kafka persists all data to disk, essentially all writes go to the **page cache** of OS, i.e. RAM.
- Fast **reads**:
  - Very efficient to transfer data from page cache to a network **socket**
  - Linux: **sendfile()** system call
- Combination of the above two features → highly efficient Kafka



---

**PROCESSES: PRODUCERS**

A solid green horizontal bar at the bottom of the slide.

# Producer intro



- Producer **use cases**:
  - Record user activities for auditing and analysis
  - Record infrastructure metrics, e.g. CPU load, RAM use, bandwidth utilization
  - Store logs
  - Record information from smart devices, e.g. in an electric power system setting
  - Buffer information before writing to a database
- Producers create **producer record** objects which consist of (topic, partition, key, value) tuples
  - The partition and key values are optional
  - When defined, the partition defines the destination partition → if it undefined, the partitioner assigns the record
- Producers rely on different serializers to convert records, e.g. Apache Avro, Java serialization, custom serialization

# Producer workflow

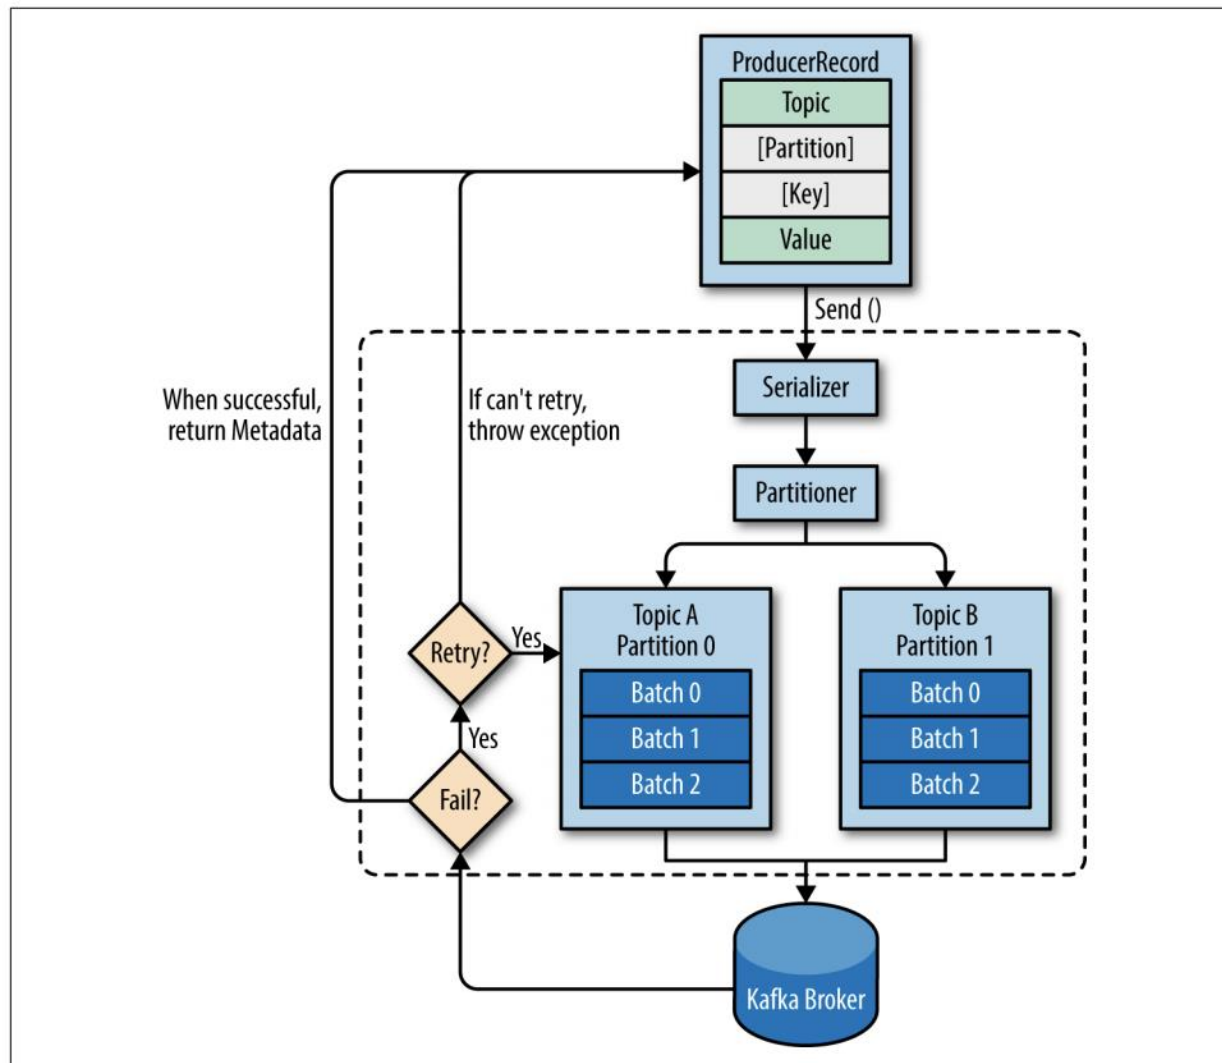


Figure 3-1. High-level overview of Kafka producer components

---

**PROCESSES: CONSUMERS**

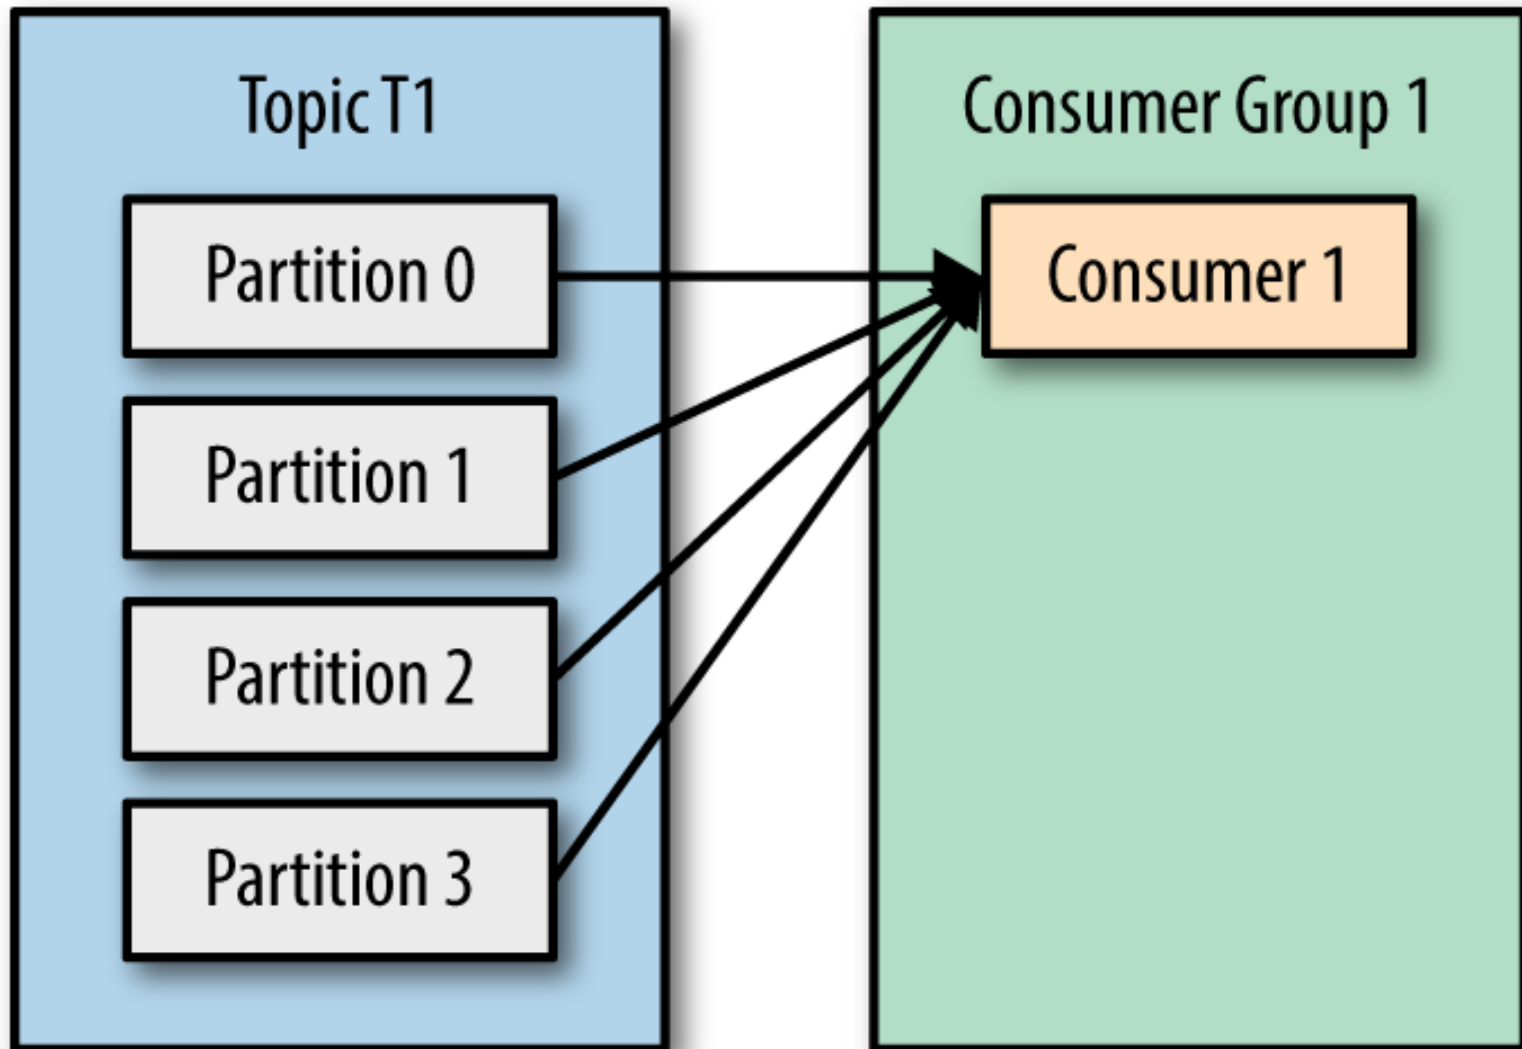
A solid green horizontal bar at the bottom of the slide.

# Consumer intro

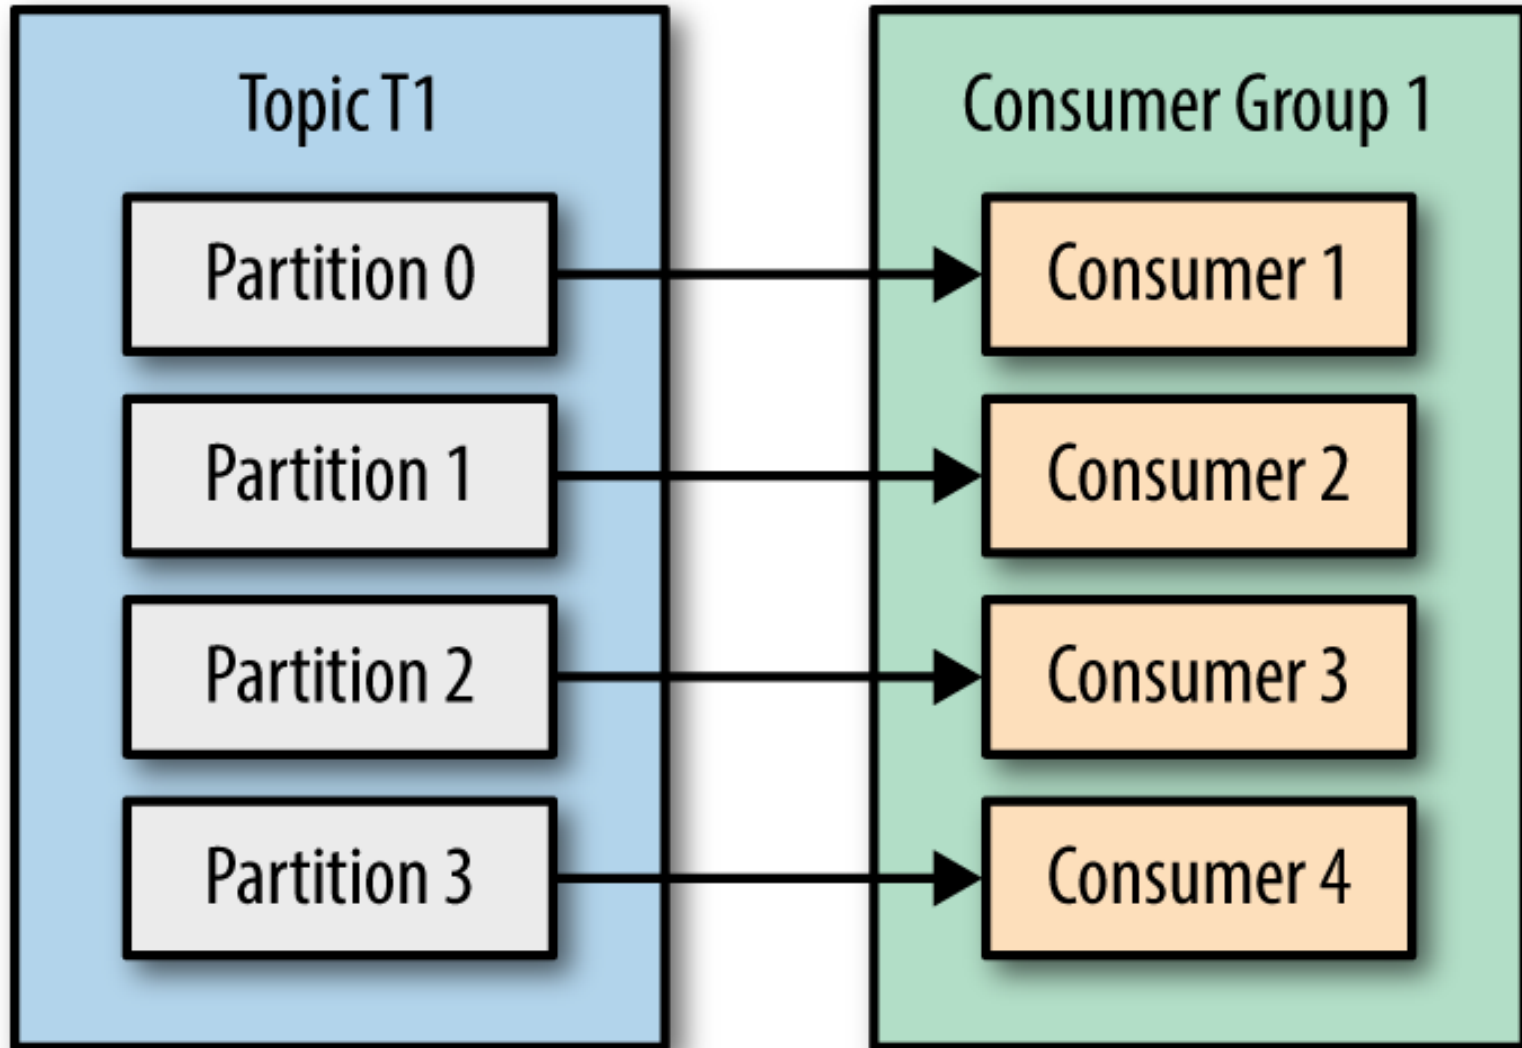


- **DEF: Kafka consumers** are processes which subscribe to and receive messages from Kafka topics
- **Simple consumer** scenario: a single consumer subscribe to a single or multiple topics and processes the data
  - A single consumer can become a bottleneck if it performs costly data analyses or writes to a database → multi-consumer usage scenario
- **Multi-consumer** scenario: a single consumer process cannot process the tide of incoming, unbounded data flows → consumer groups with multiple consumers
  - The different consumers receive messages from a different subset of topic partitions
- When consumers consume messages, they commit the current partition offsets to a special Kafka topic
  - Earlier versions (prior to 0.10.x) committed offsets to Zookeeper

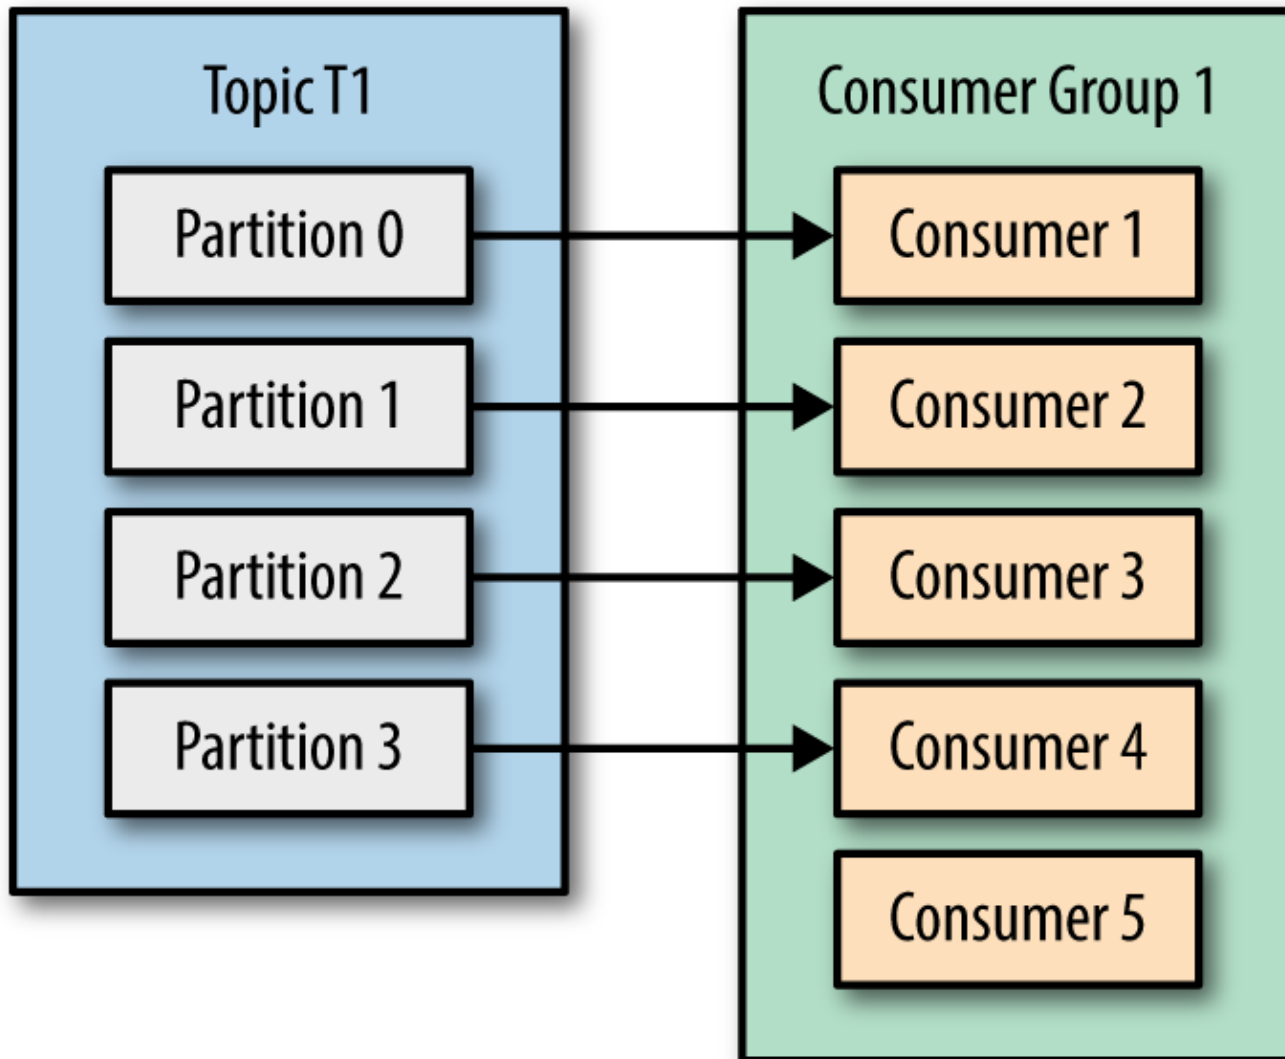
# Single group, single consumer



# Optimum number of consumers

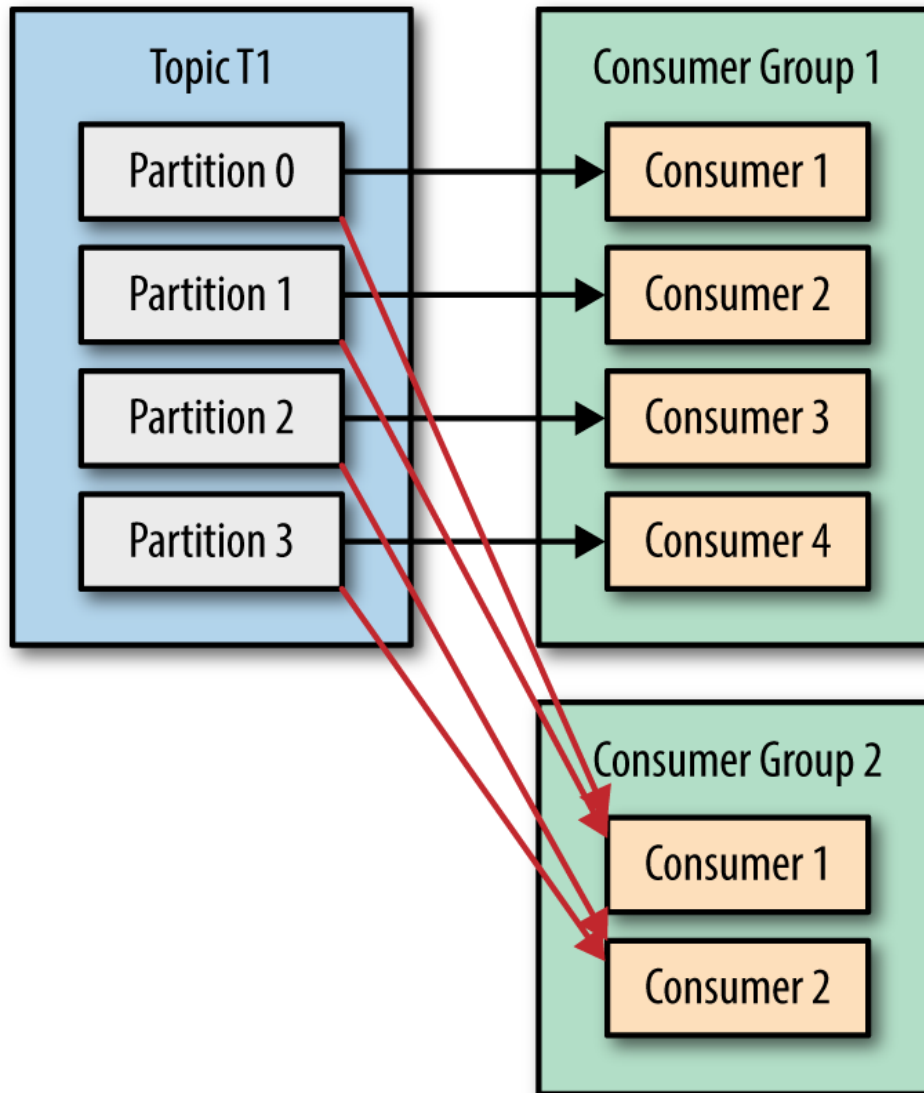


# Too many consumers





# Multiple consumer groups



- The different consumer groups can perform different types of data transformations
- Note: adding additional consumers on-the-fly is possible → partition rebalance

---

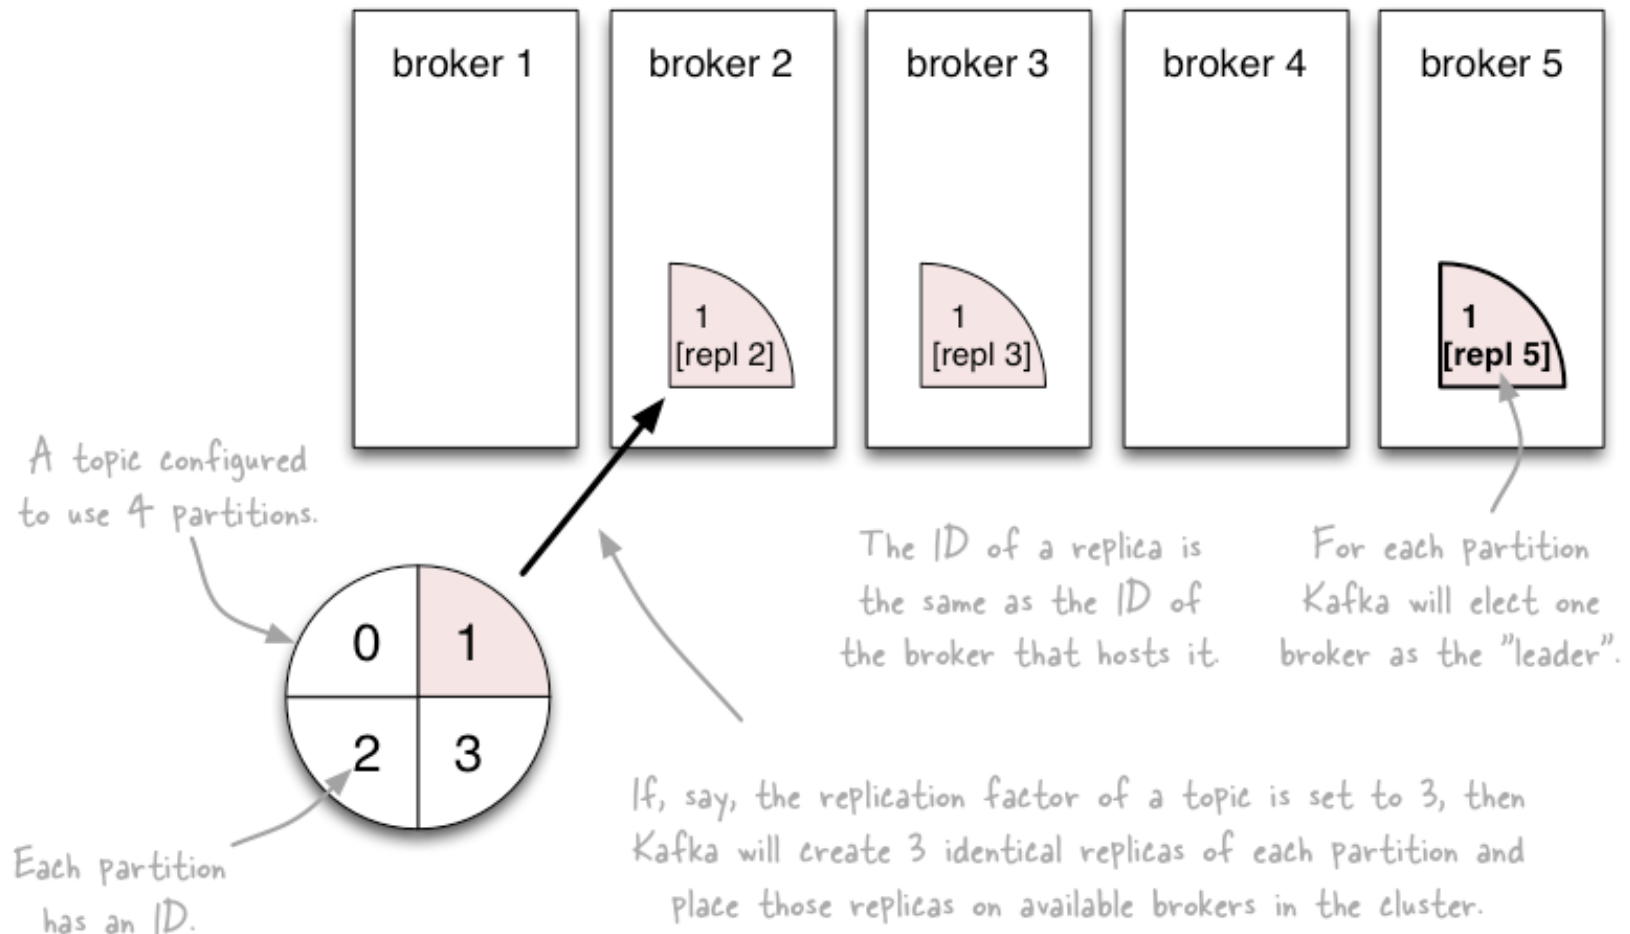
# REPLICATION AND CONSISTENCY

# Replicas of a partition

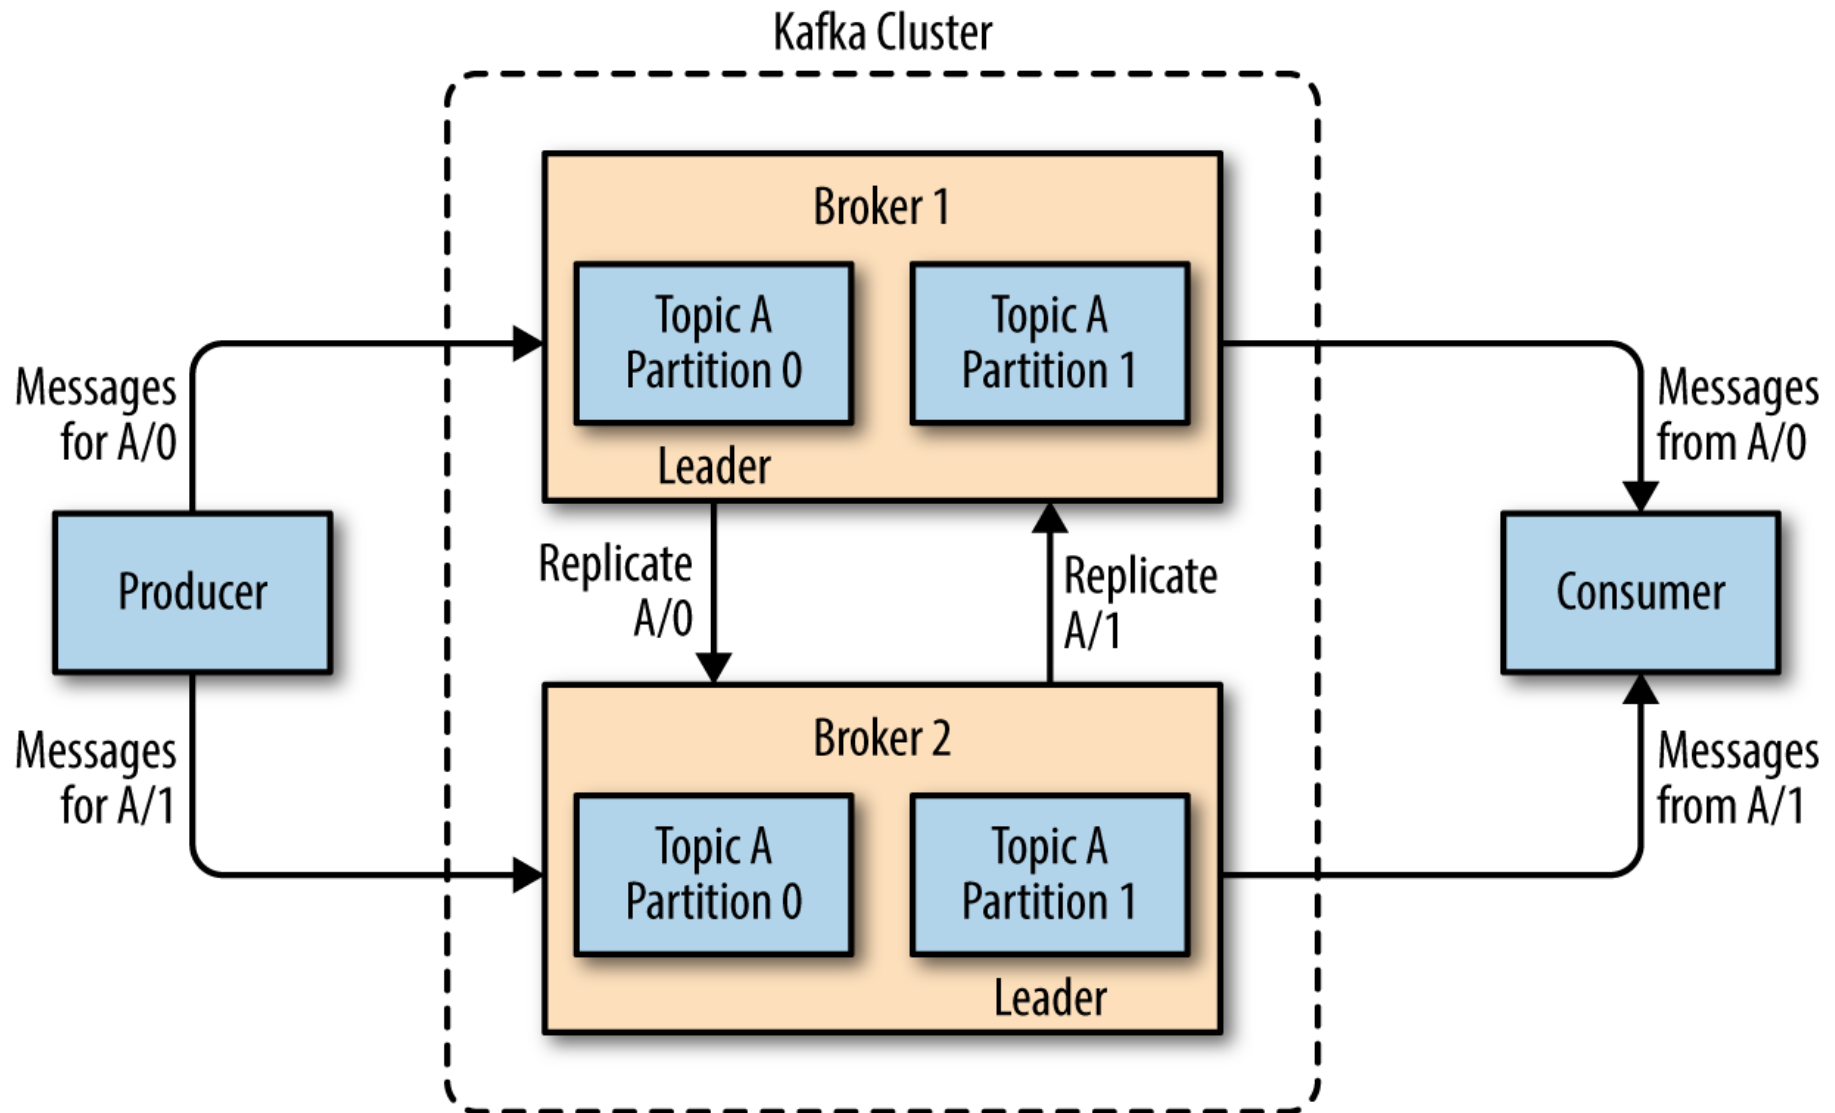


- Kafka **replicas** are partition “backups”
  - Kafka tolerates  $(numReplicas - 1)$  dead brokers before losing data
- Replica types:
  - **Leader replica:** Each partition has a single replica designated as the leader. All produce/consume requests go to the leader.
  - **Follower replica:** All non-leader replicas for each partition are called followers. Followers do not serve produce/consumer requests. They exist to avoid data loss.
- It is the task of the leader to know which follower replicas are up-to-date → in-sync replicas
- Only in-sync replicas can become leaders when the current leader exits the Kafka cluster, e.g. it fails

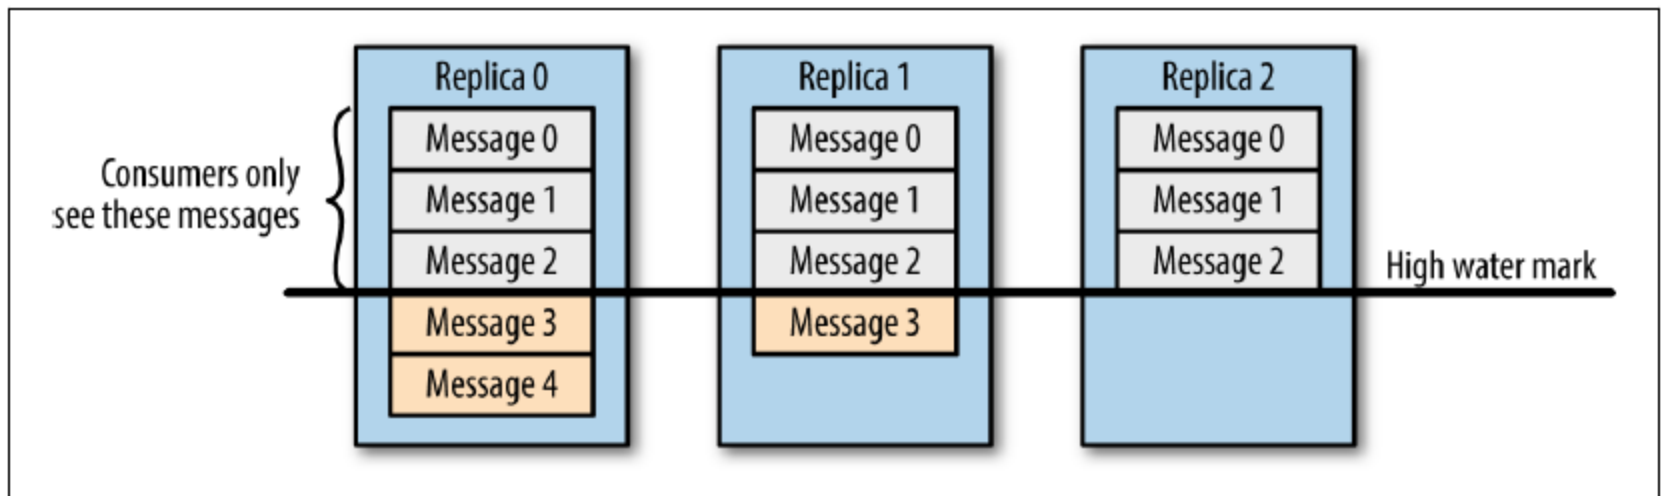
# Topics vs. Partitions vs. Replicas



# Partition replication



# Replication and visibility



*Figure 5-4. Consumers only see messages that were replicated to in-sync replicas*

---

# **MONITORING AND CONTROL**



# Monitoring Kafka

- Nothing fancy built into but see:
  - <https://cwiki.apache.org/confluence/display/KAFKA/System+Tools>
  - <https://cwiki.apache.org/confluence/display/KAFKA/Ecosystem>

Kafka Web Console


Zookeepers

Brokers

Topics

Zookeeper	Topic	Partitions
Test SNAPSHOT	dead-letter-queue	1
Latest	exceptions	1
Latest	metrics	10
Latest	alerts	2
Latest	logs	5
Latest	notifications	10
Live	exceptions	1
Live	metrics	10
Live	alerts	2
Live	logs	5
Live	notifications	10

Kafka Web Console



Kafka Offset Monitor



---

# SUMMARY



# Who uses Kafka?

- **LinkedIn:** (user) activity streams, operational metrics, data bus
  - 400 nodes, 18k topics, 220B msg/day (peak 3.2M msg/s), May 2014
- **Netflix:** real-time monitoring and event processing
- **Twitter:** as part of their Storm real-time data pipelines
- **Spotify:** log delivery (from 4h down to 10s), Hadoop
- **Loggly:** log collection and processing
- **Mozilla:** telemetry data
- Others: Airbnb, Cisco, Gnip, InfoChimps, Ooyala, Square, Uber, etc.

# Summary



- Introduction
- Motivation
- Architecture
- Processes
  - Brokers
  - Producers
  - Consumers
- Replication & consistency
- Monitoring & control



---



**Thank you for your attention!**