



WATERMARKS AND WINDOWING (W2)

Stream mining (SM)

Imre Lendák, PhD, Associate Professor

Péter Kiss, PhD candidate

Outline

- Advanced watermarking
- Advanced windowing



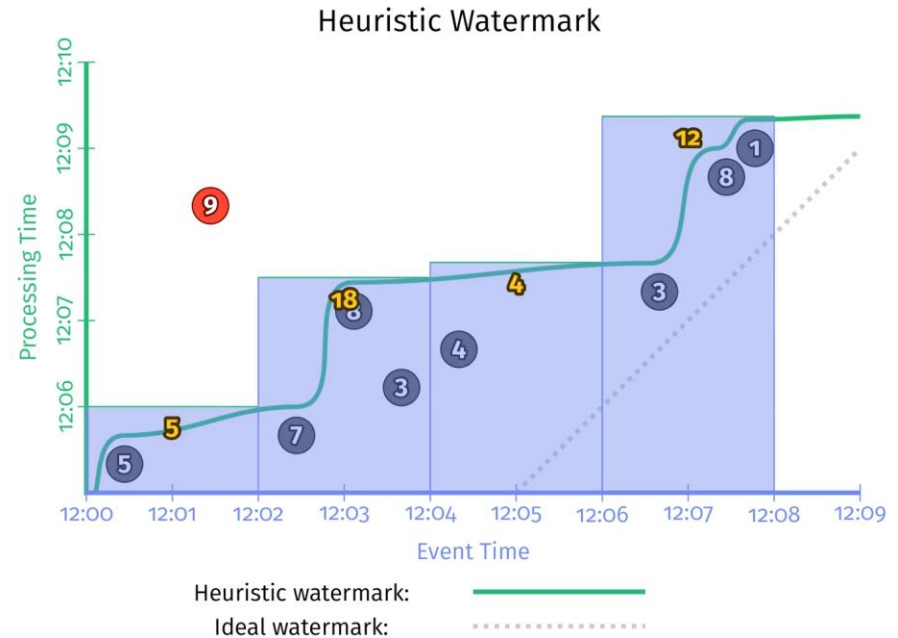

Stream mining

ADVANCED WATERMARKING

Watermarks

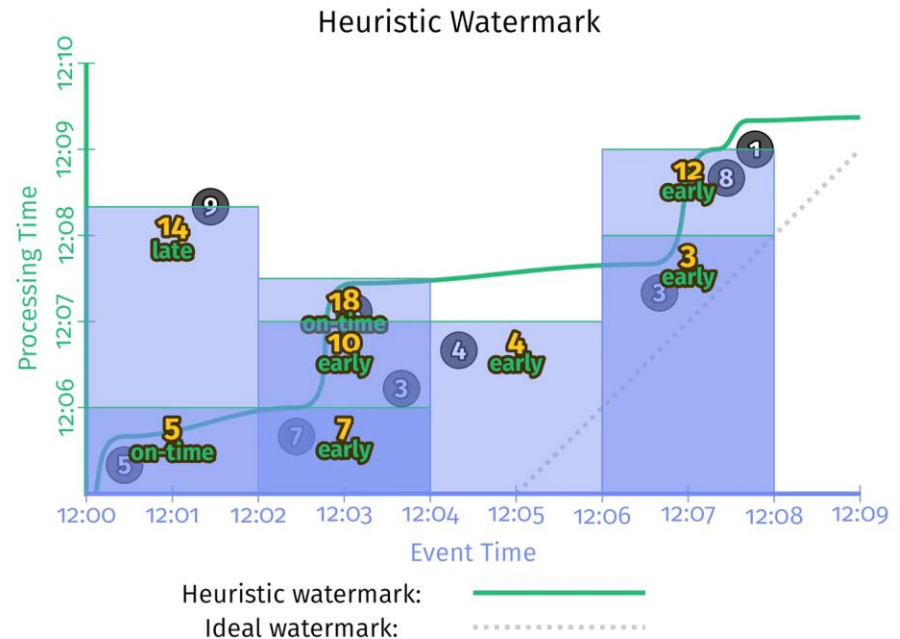

- **DEF #1:** Watermarks are temporal notions of input completeness in the event time domain
 - ‘Guesses’ about data completeness, i.e. all relevant data received
- **DEF #2:** Watermarks allow streaming systems to measure progress and completeness relative to event times of the records being processed
- Watermark varieties:
 - Perfect watermark = usable when we have perfect knowledge about the input data → all data on time (i.e. no late data)
 - Heuristic watermark = provide an estimate of (data) progress as good as it gets → usually based on domain knowledge
- **NOTE:** Additional watermark management challenges in data processing pipelines with multiple processing steps

Perfect vs heuristic watermarks




<http://streamingsystems.net/fig/2-10>

Watermarks meet triggers



- Early/on-time/late triggers
 - Zero or more early triggers **periodically** firing
 - Single **on-time** trigger on completeness/watermark
 - Zero or more triggers for **late data** unaccounted for by the watermark

Allowed lateness



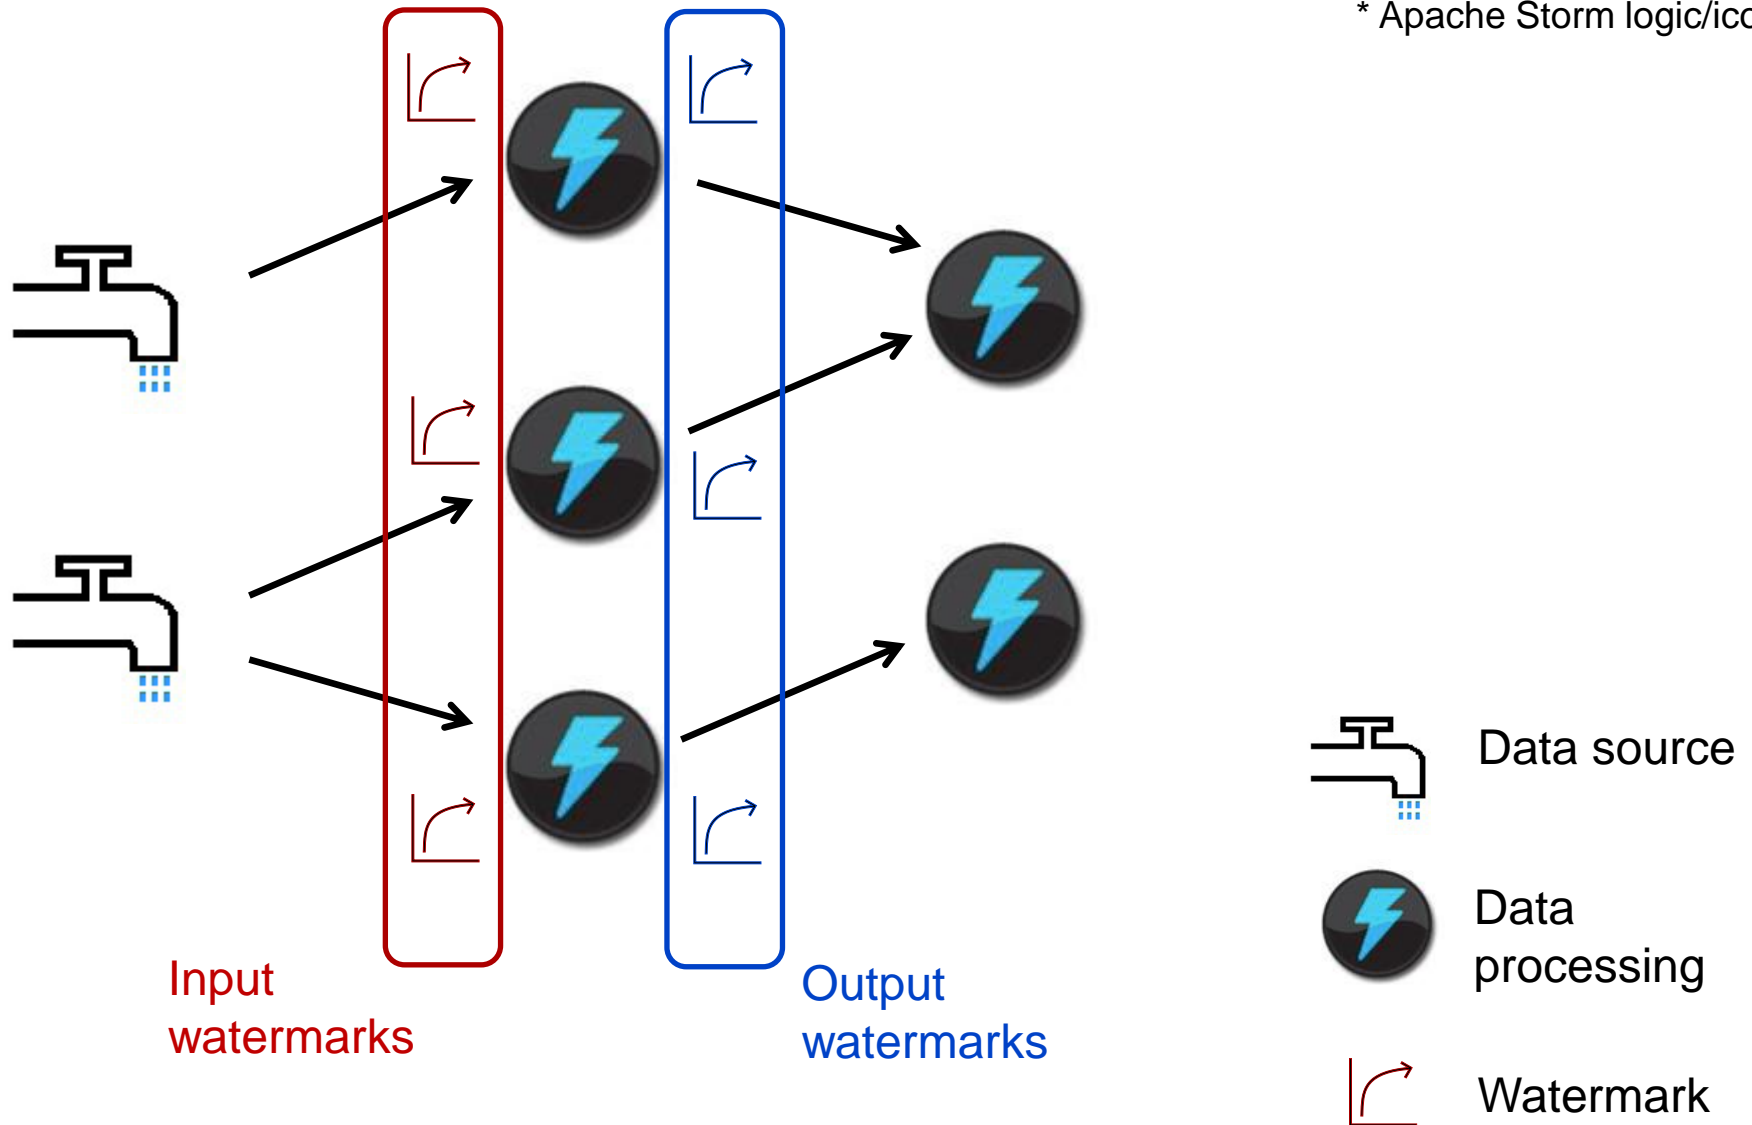
Heuristic watermark: —————

Ideal watermark:

<http://streamingsystems.net/fig/2-12>

Watermark @ boundaries

* Apache Storm logic/icons



Input watermarks

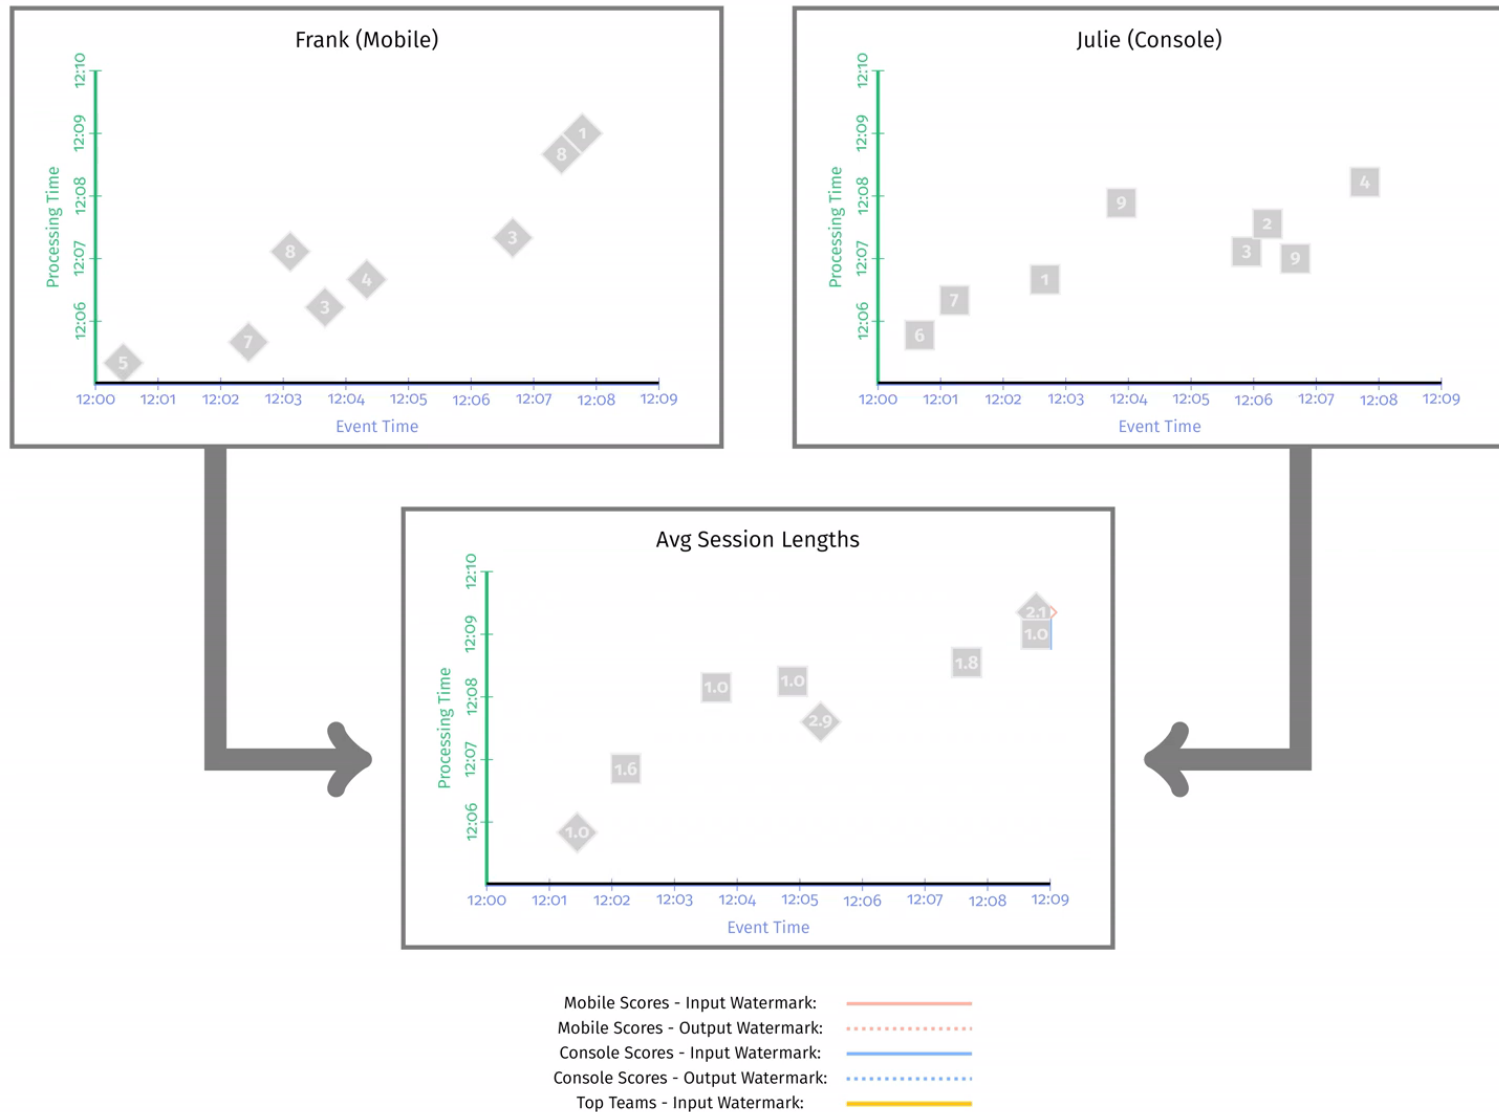
- **Input watermark** = calculated based on the data processing pipeline's progress upstream ('to the left') of the input of the current processing stage
- For (data) sources there is a specific function generating the watermark. It is created based on our understanding of the input data source
- For non-sources it is the **minimum** of all dependent output watermarks upstream
 - Why the minimum?

Output watermarks

- Each processing node can **buffer** active messages until some operation is completed, e.g. aggregation in a previous stage.
 - Each buffer can be tracked with its own watermark
- **DEF:** Output watermarks are calculated at the output(s) of processing nodes and incorporate/observe input + processing (time) delay
- Output watermarks are (usually) based on a combination of the following watermarks
 - Data source → source watermarks
 - External input → other external sources, e.g. upstream processing
 - (Internal) state → there is an internal state-machine inside the processing node with clearly defined rules of progress
 - Per-output buffer → times of (data) processing results written to output buffers




Watermark propagation



<http://streamingsystems.net/fig/3-5>

Percentile watermarks




100% Watermark:
66% Watermark:
33% Watermark:
Ideal Watermark:



<http://streamingsystems.net/fig/3-11>

Processing time watermarks



- The processing time watermark is constructed based on the timestamp of the oldest (processing) operation not yet completed
 - Allows insight into processing delay separate from data delay
 - Growing processing watermark → faults preventing (process) operations to complete → user/administrator action needed
- Useful to distinguish (streaming) system latency from data (source) latency


Stream mining

WINDOWING

Advanced windowing intro

- Processing-time windowing is useful for use cases in which data is observed as received, e.g. web server traffic
 - Implementation via triggers or ingress time
 - Time-sensitive → use only when data is received in perfect order or when we do not care about when events actually happened
- Event-time windowing is used in use cases in which the exact sequence of events is relevant, e.g. billing, user behavior
 - Order agnostic, i.e. does not care about the order in which the data is received as these solutions are able to re-order pieces of data received from different sources in jumbled order

General windowing strategies




<http://streamingsystems.net/fig/1-8>

- Fixed windows → Slice time into fixed-sized temporal length, usually across all keys/variables.
- Sliding windows → Defined by fixed length & fixed period.
- Sessions → Sequences of events terminated by a gap.

Stream mining

PROCESSING-TIME WINDOWING

Windowing by processing time




<http://streamingsystems.net/fig/1-9>

- Windows are created based on absolute (processing) time, i.e. data is put into windows based on the order they arrive in

Processing-time windowing & triggers

- Triggers fire periodically in the processing-time domain and initiate data processing, e.g. every 2 minutes
- Accumulating or discarding depending on the use case
- Most traditional streaming systems worked like this

PT with triggers in action



- Processing-time windowing via triggers can create different results for the same input data set with different processing-time ordering

Processing-time windowing & ingress time

- Ingress time windowing basics
 - Event times are replaced with ingress times, i.e. the times when they were received by the streaming system
 - Utilize perfect watermarks made possible via ingress time → processing is triggered when the watermark is reached
 - Accumulation mode is based on the use case
- Outputs can be different for different ordering of the same set of input data

PT with ingress time in action




- Note: a perfect watermark is made possible via the use of ingress time

Stream mining

EVENT-TIME WINDOWING


Windowing by event time – Fixed




<http://streamingsystems.net/fig/1-10>

- Data is collected and windowed based on times at which it occurred, i.e. event time.

Event-time windowing in action



Heuristic watermark: —
Ideal watermark:



Heuristic watermark: —
Ideal watermark:

- Sum calculus in event-time windows over the same pieces of data in two different processing time orderings.


Variation #1: Aligned fixed windows

- Aligned fixed windows start and end at the same time for all keys/observations
 - Example: thousands of measurement points on a factory floor with advanced motor monitoring with high frequency data sampling for early fault detection, which are processed every minute (in absolute, processing time)
- Pro:
 - Implementing fixed windows is relatively easy
- Contra:
 - The data is seldom occurring in a windowed fashion in real-life streaming systems
 - There is a high processing (CPU, RAM) load at the end of the windows when all windows are processed at the same time

Variation #2: Unaligned fixed windows

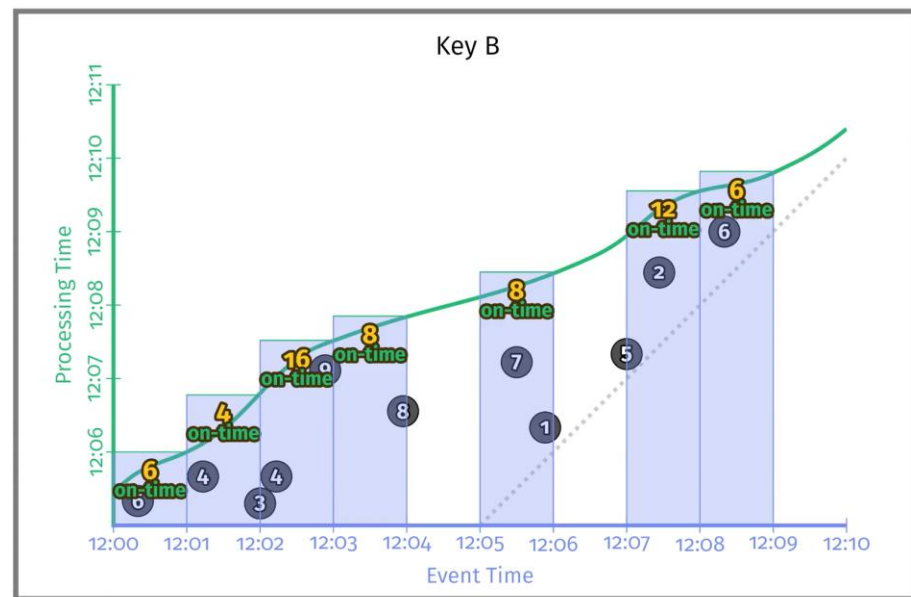
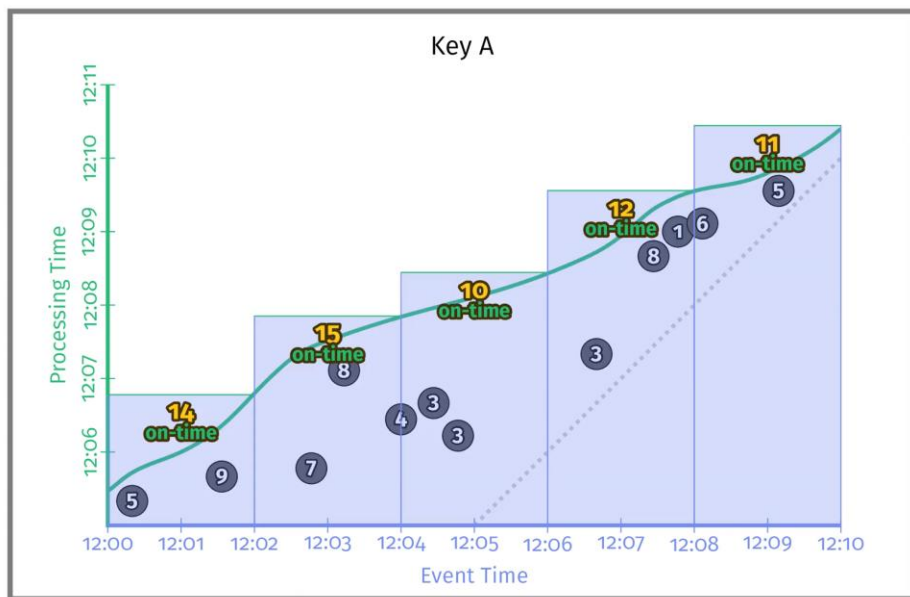
- Unaligned fixed windows do not start at the same time for all keys/observations
- The length of these windows is still pre-defined
- Pro:
 - More evenly spread processing loads
- Contra:
 - The data is seldom occurring in a windowed fashion in real-life streaming systems
 - Not as easy to implement

Unaligned fixed windows in action



- The horizontal lines mark the 'ends' of the windows → that's when processing is triggered

Variation #3: Per-element/key fixed windows




- Fixed windows with different key sizes:
 - Key A with fixed window size = 2 minute(s)
 - Key B with fixed window size = 1 minute(s)
- Useful when different stream analysis requirements exist

Stream mining

SESSIONS IN WINDOWING

Windowing by event time – Sessions



<http://streamingsystems.net/fig/1-11>


- Data is collected into sessions based on temporal proximity and key/user.

Session windows


- Session windows are data-driven → the location and length of windows is dictated by the input data itself
 - This is opposed to fixed windows, which are finalized after the expiration of a certain time period either in event or processing time
- Sessions are unaligned
 - There is absolutely no guarantee that data will occur at the exact same time for different keys
- Session creation can be based on
 - Predefined (time) gap which separates two sessions – this is done more often
 - Predefined tag which is assigned to each piece of data – this is a less frequent scenario

Proto sessions & merging


Unmerged Proto-sessions - 60 min each



Merged Session #1 - 135 min



Merged Session #2 - 98 min

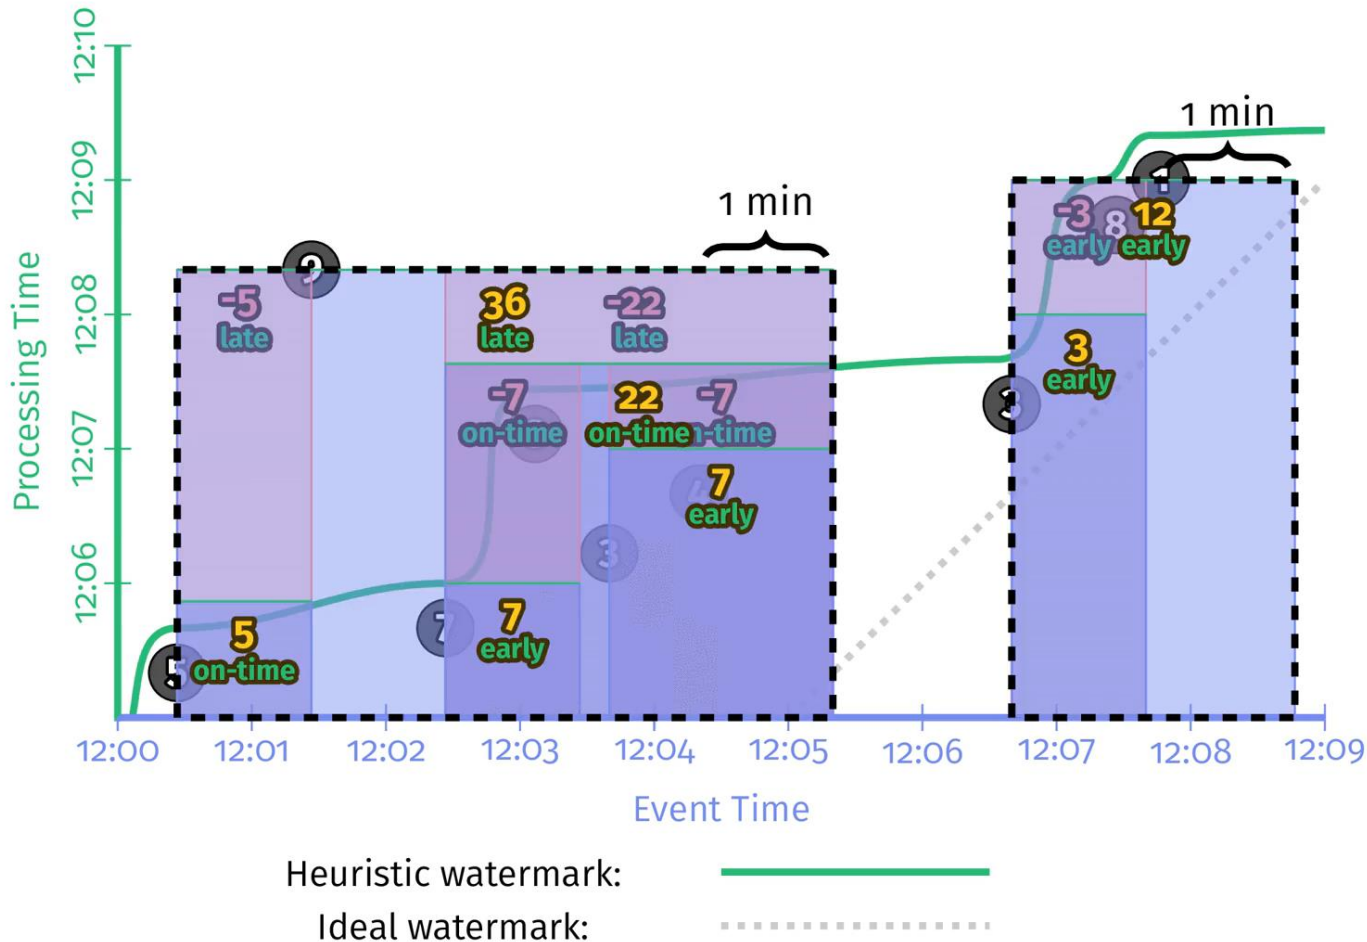


<http://streamingsystems.net/fig>

Session creation steps

- Phase 0: The gap duration is defined as the length of time between two sessions
- Phase I: Assignment
 - Each element is initially placed in a proto-session window
 - The proto-session starts with the occurrence of the event
 - The proto-session extends for the gap duration
- Phase II: Merging
 - A grouping strategy is defined
 - All eligible proto-sessions are sorted
 - All overlapping proto-sessions are grouped into sessions

Session creation in action

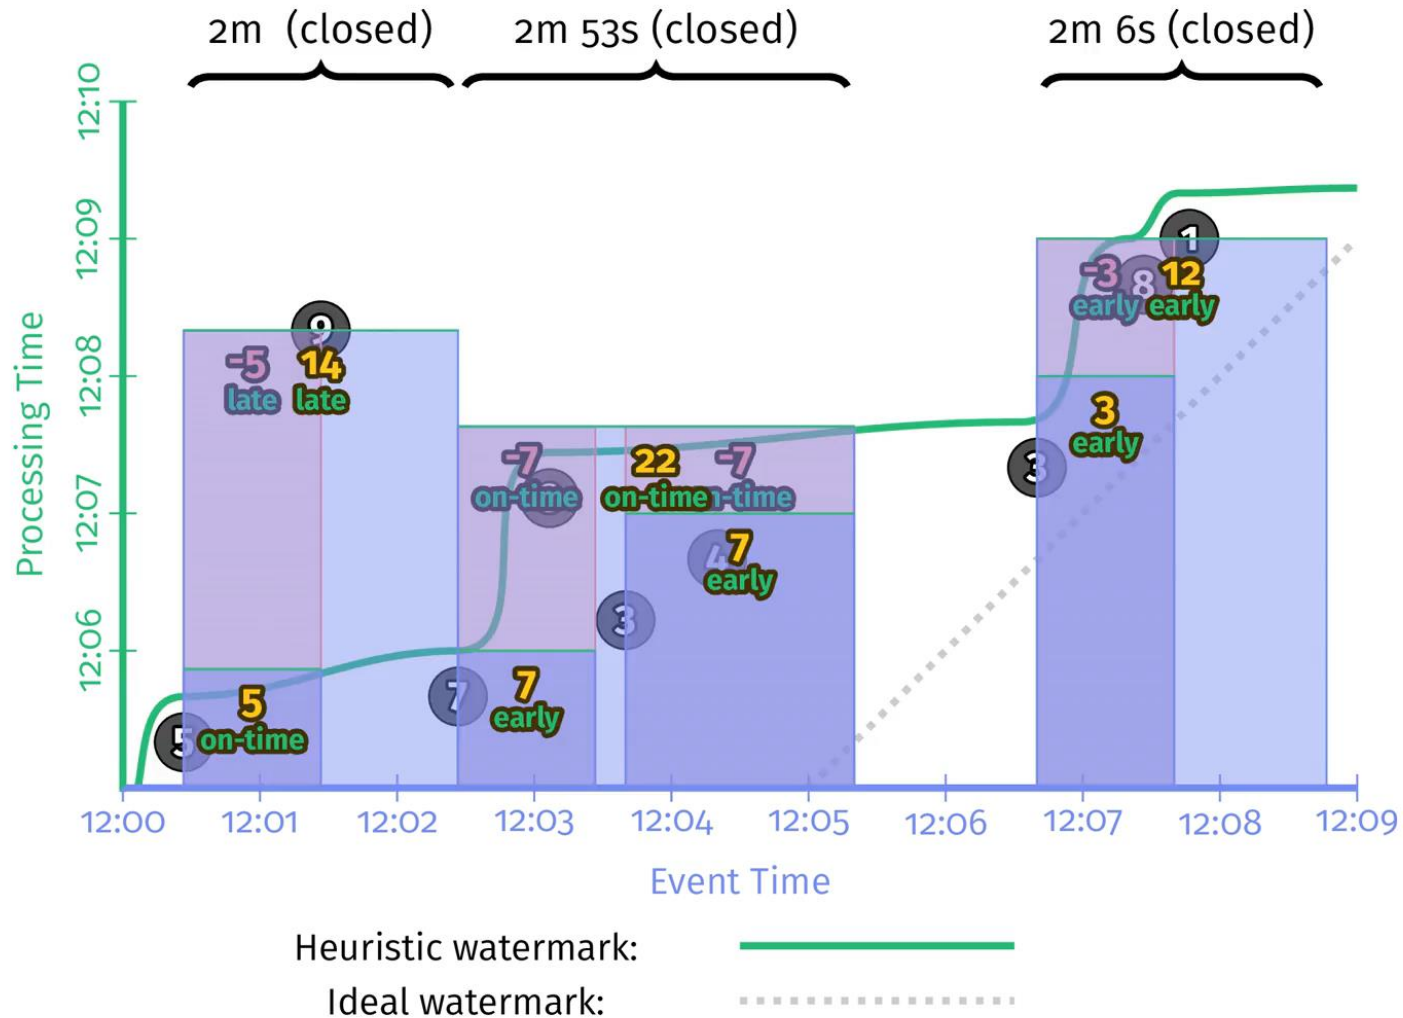


<http://streamingsystems.net/fig>

Variation #1: Bounded sessions

- Bounded sessions are not allowed to grow beyond a predefined size
- The 'size' can be defined in time, element count, aggregate value or other dimension
- Example: time-bounded sessions, e.g. a maximum session length of 3 minutes

Bounded sessions in action



<http://streamingsystems.net/fig>

Summary

- Watermarks
 - Perfect vs heuristic
 - Input & output
 - Percentile
 - Processing time
- Advanced windowing
 - Processing-time
 - Event-time
 - Sessions
 - No one-size-fits-all windowing strategy !



Thank you for your attention!